

Chapter 1: Introduction to Rancher and Kubernetes

No Images...

Chapter 2: Rancher and Kubernetes High-Level Architecture

```
! cluster.yaml
1  nodes:
2    - address: 172.27.7.21
3      user: root
4      hostname_override: alubrancherl01
5      internal_address: 172.27.7.21
6      role: [controlplane,worker,etcd]
7    - address: 172.27.7.22
8      user: root
9      hostname_override: alubrancherl02
10     internal_address: 172.27.7.22
11     role: [controlplane,worker,etcd]
12    - address: 172.27.7.23
13      user: root
14      hostname_override: alubrancherl03
15      internal_address: 172.27.7.23
16      role: [controlplane,worker,etcd]
17
18  services:
19    etcd:
20      backup_config:
21        enabled: true      # enables recurring etcd snapshots
22        interval_hours: 12 # time increment between snapshots
23        retention: 6       # time in days before snapshot purge
24      s3backupconfig:
25        access_key: "ABCDEFGHijklmnpO..."
26        secret_key: "123456789abcdefghijklmpo..."
27        bucket_name: "etcd-backups"
28        folder: "Cluster-Name-Here"
29        endpoint: "s3.us-west-1.wasabisys.com"
30        region: "us-west-1"
31
32  dns:
33    provider: coredns
34    upstreamnameservers:
35      - 172.27.2.23
36      - 172.27.2.24
```

Chapter 3: Creating a Single Node Rancher

```
root@a1ubranl00:~# docker info
Client:
 Context:    default
 Debug Mode: false
 Plugins:
  app: Docker App (Docker Inc., v0.9.1-beta3)
  buildx: Build with BuildKit (Docker Inc., v0.6.1-docker)

Server:
 Containers: 0
  Running: 0
  Paused: 0
  Stopped: 0
 Images: 0
 Server Version: 20.10.8
 Storage Driver: overlay2
  Backing Filesystem: extfs
  Supports d_type: true
  Native Overlay Diff: true
  userxattr: false
 Logging Driver: json-file
 Cgroup Driver: cgroupfs
 Cgroup Version: 1
 Plugins:
  Volume: local
  Network: bridge host ipvlan macvlan null overlay
  Log: awslogs fluentd gcplogs gelf journald json-file local logentries splunk syslog
 Swarm: inactive
 Runtimes: io.containerd.runc.v2 io.containerd.runtime.v1.linux runc
 Default Runtime: runc
 Init Binary: docker-init
 containerd version: e25210fe30a0a703442421b0f60afac609f950a3
 runc version: v1.0.1-0-g4144b63
 init version: de40ad0
 Security Options:
  apparmor
  seccomp
   Profile: default
 Kernel Version: 5.4.0-86-generic
 Operating System: Ubuntu 20.04.3 LTS
 OSType: linux
 Architecture: x86_64
 CPUs: 4
 Total Memory: 7.748GiB
 Name: a1ubranl00
 ID: KRPX:05TH:3XMS:GS5L:UPP3:CEBZ:X25W:QCTA:RBJN:IG76:6L2U:JZC2
 Docker Root Dir: /var/lib/docker
 Debug Mode: false
 Registry: https://index.docker.io/v1/
 Labels:
 Experimental: false
 Insecure Registries:
  127.0.0.0/8
 Live Restore Enabled: false

WARNING: No swap limit support
root@a1ubranl00:~#
```

```
root@a1ubranl00:~# cat /etc/docker/daemon.json
{
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "10m",
    "max-file": "3"
  }
}
root@a1ubranl00:~# █
```

```
-----BEGIN CERTIFICATE-----
Root certificate
-----END CERTIFICATE-----
-----BEGIN CERTIFICATE-----
Intermediate certificate
-----END CERTIFICATE-----
-----BEGIN CERTIFICATE-----
Server certificate
-----END CERTIFICATE-----
```

```
-----BEGIN RSA PRIVATE KEY-----
Proc-Type: 4,ENCRYPTED DEK-Info: DES-EDE3-CBC,
....
-----END RSA PRIVATE KEY-----
```

```
-----BEGIN ENCRYPTED PRIVATE KEY-----
MIIJnDB0BgkqhkiG9w0BBQ0wQTApBgkqhkiG9w0BBQwwHAQIupZ5LBxkx4wCAggA
....
0+S3p1U3GAYMxdbZAcMtKnQzSVI4AalbF7a+7C1bFS4JIWBg/W1jkzZP/lc6klKq
b1J+hEz5rSoD+E/2ccsLpg==
-----END ENCRYPTED PRIVATE KEY-----
```

```
-----BEGIN RSA PRIVATE KEY-----
MIIJKQIBAAKCAgEAuozVjS468biRJAwrR0+LLW+fxoucW4u5vsI4UyFnLA2KnNMV
vc8idiTLTy0jukxgUYCABGxu0jzk5QP9Fgqs7p2+MmZn+lDBPF9zLAP/SJQVL2Jx
....
Br8vCDapjAgW2pazmpzDCv67C3G6y04WTqBAphSpk4AIy7YJUgDQxnf3sfMgCiqV
ZIPf77ywmS0F50u6EyTYqglSGRXYXfUrR1j+HmqQ+EdMNEY1Kk9lwklvIZ/p
-----END RSA PRIVATE KEY-----
```

```
docker run -d \
--name rancher_server \
--restart=unless-stopped \
-p 80:80 \
-p 443:443 \
-v /etc/rancher/ssl/tls.crt:/etc/rancher/ssl/cert.pem \
-v /etc/rancher/ssl/tls.key:/etc/rancher/ssl/key.pem \
--privileged \
rancher/rancher:v2.5.8 \
--no-cacerts
```

```
root@a1ubranl00:~# docker logs rancher_server 2>&1 | grep "Bootstrap Password:"
2021/09/28 01:40:54 [INFO] Bootstrap Password: p7brsdjzs4bnfblxdn6rqhq2mfgffz5ghdjjddj7ptw5vkmjkm7q2g
root@a1ubranl00:~#
```

Welcome to Rancher!

The first order of business is to set a strong password for the default `admin` user. We suggest using this random one generated just for you, but enter your own if you like.

- ☐ Use a randomly generated password
- ☒ Set a specific password to use

New Password

••••••••

Confirm New Password

••••••••

What URL should be used for this Rancher installation? All the nodes in your clusters will need to be able to reach this.

Server URL

https://a1ubranl00.support.tools

- ☒ Allow collection of [anonymous statistics](#) to help us improve Rancher.
- ☒ I agree to the [terms and conditions](#) for using Rancher.

Continue

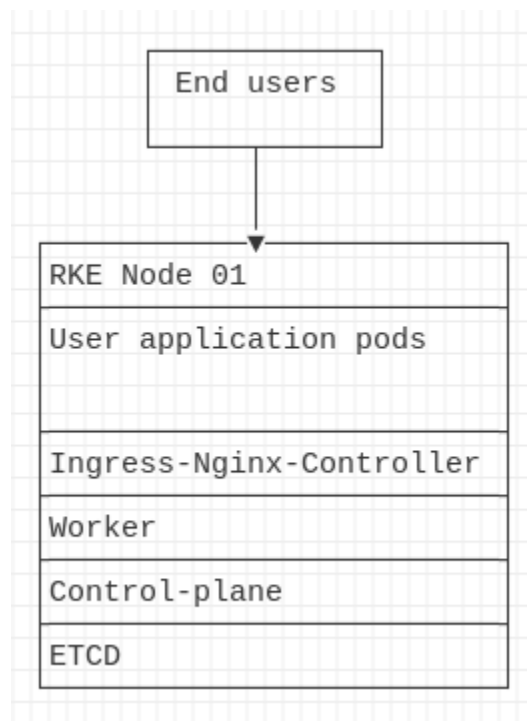
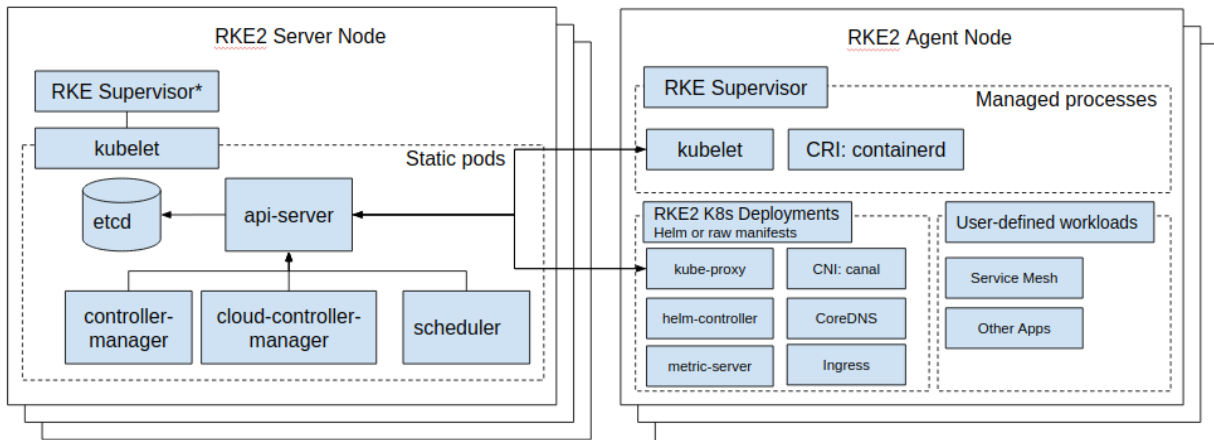
Chapter 4: Creating an RKE and RKE2 Cluster

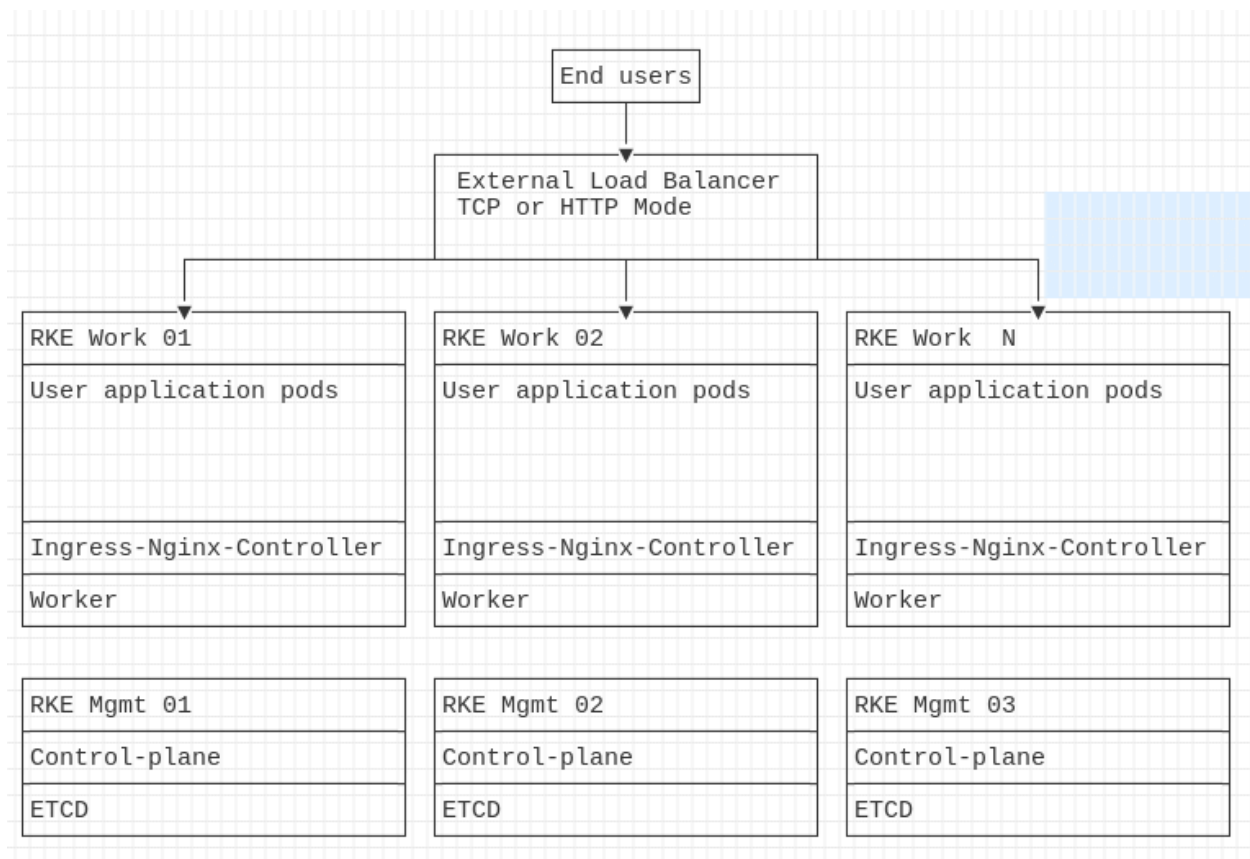
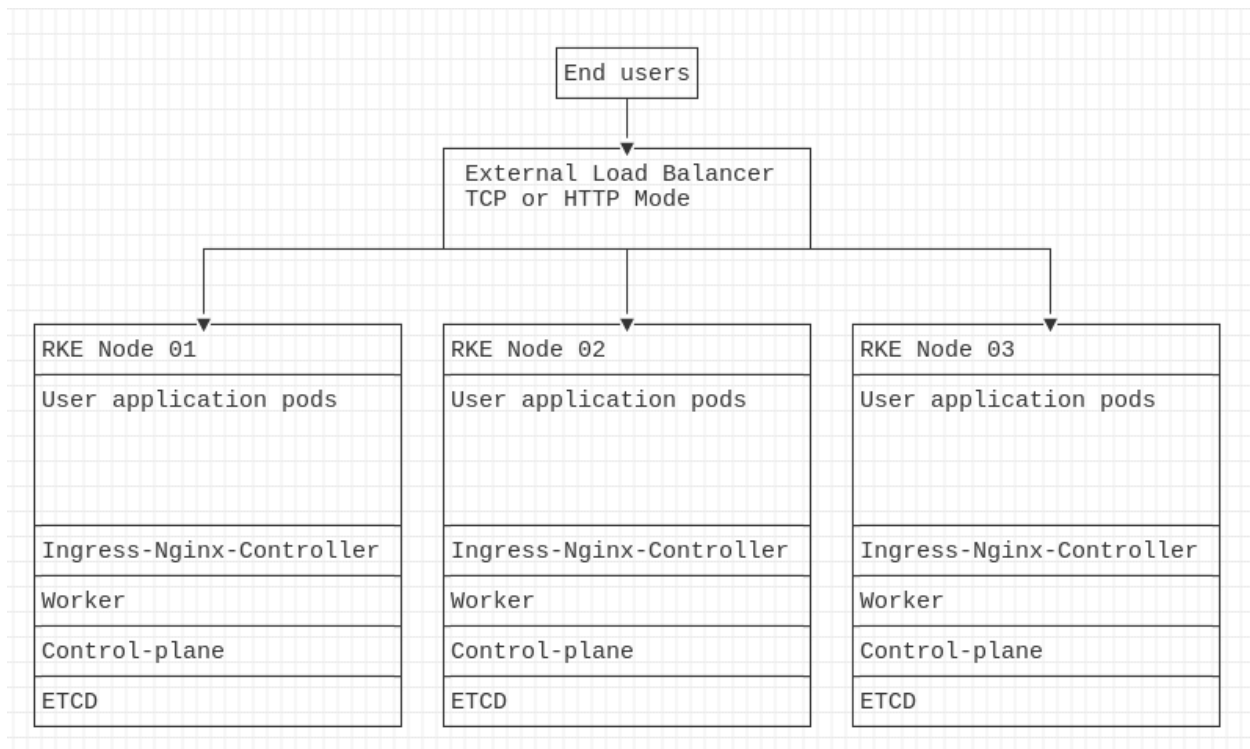
```
nodes:
- address: node01.support.tools
  hostname_override: node1
  internal_address: 192.168.1.101
  user: ubuntu
  role:
    - controlplane
    - worker
    - etcd
- address: node02.support.tools
  hostname_override: node2
  internal_address: 192.168.1.102
  user: ubuntu
  role:
    - controlplane
    - worker
    - etcd
- address: node03.support.tools
  hostname_override: node3
  internal_address: 192.168.1.103
  user: ubuntu
  role:
    - controlplane
    - worker
    - etcd

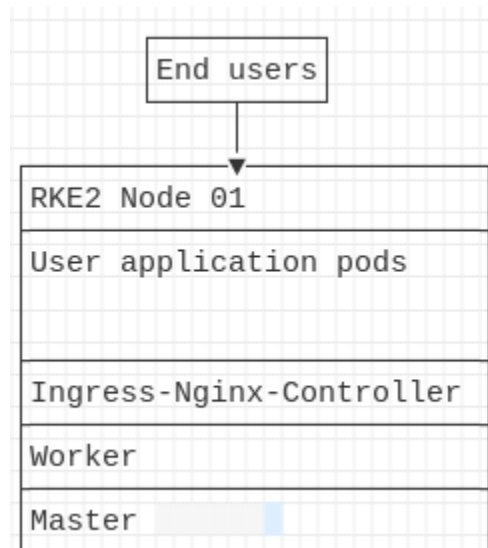
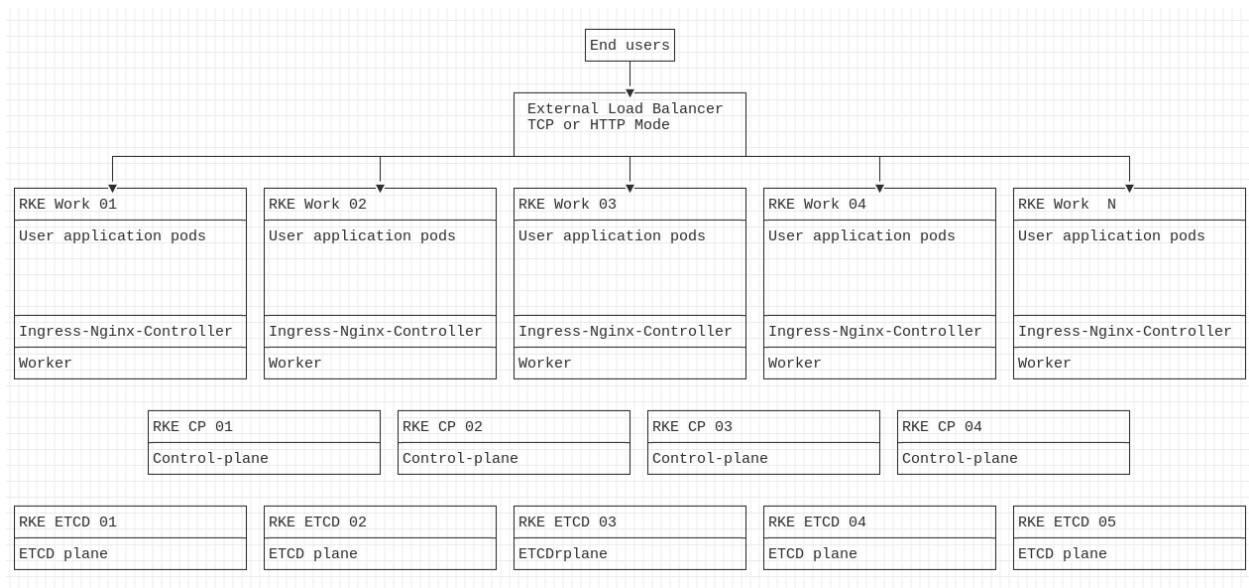
cluster_name: examplecluster
kubernetes_version: v1.21.5-rancher1-1

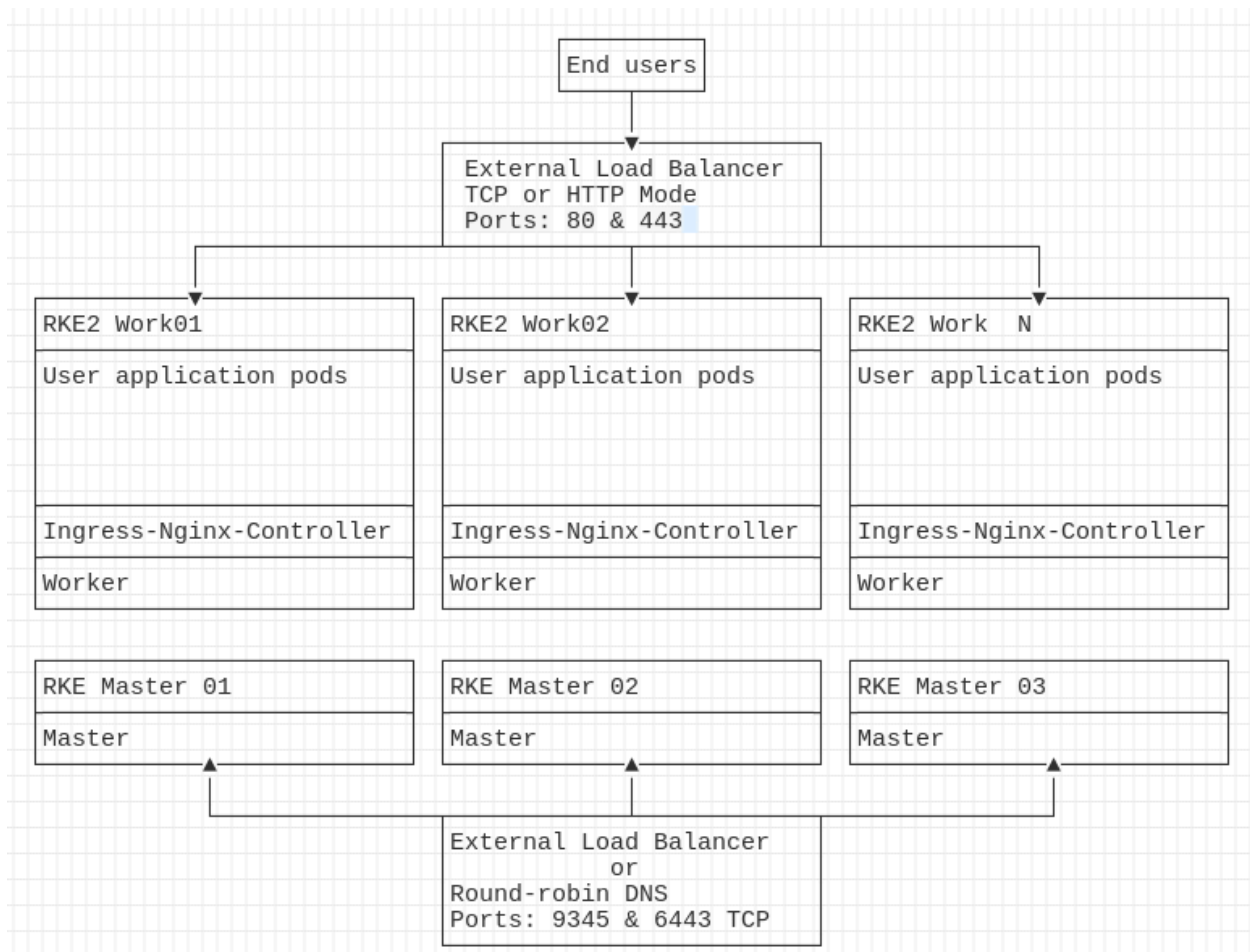
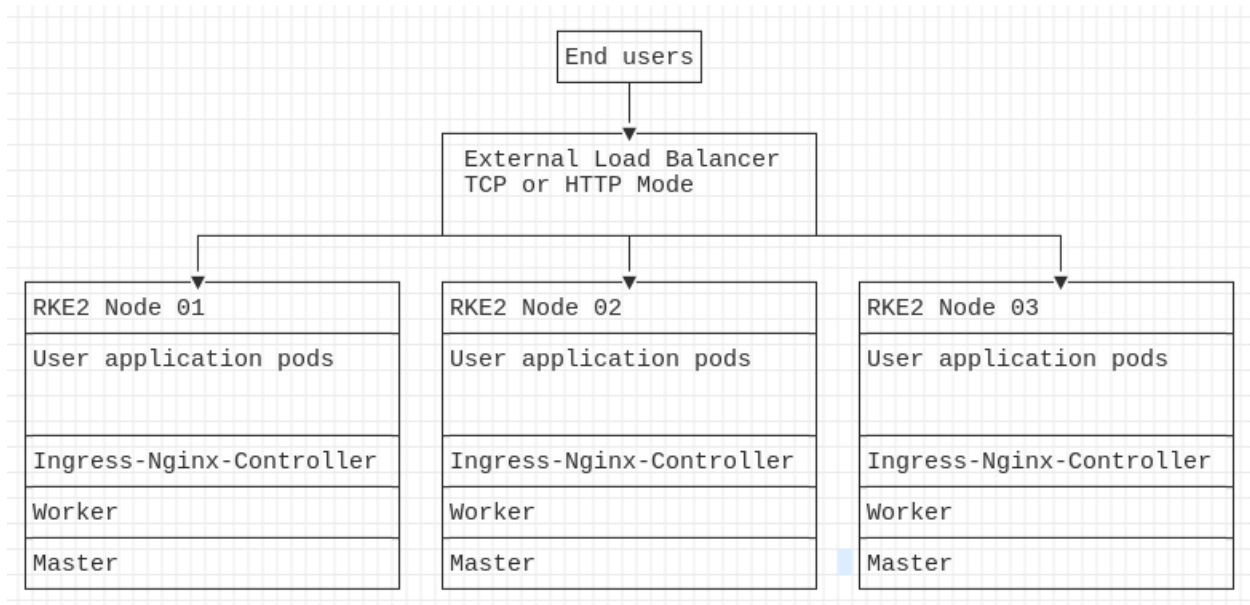
ingress:
  provider: nginx

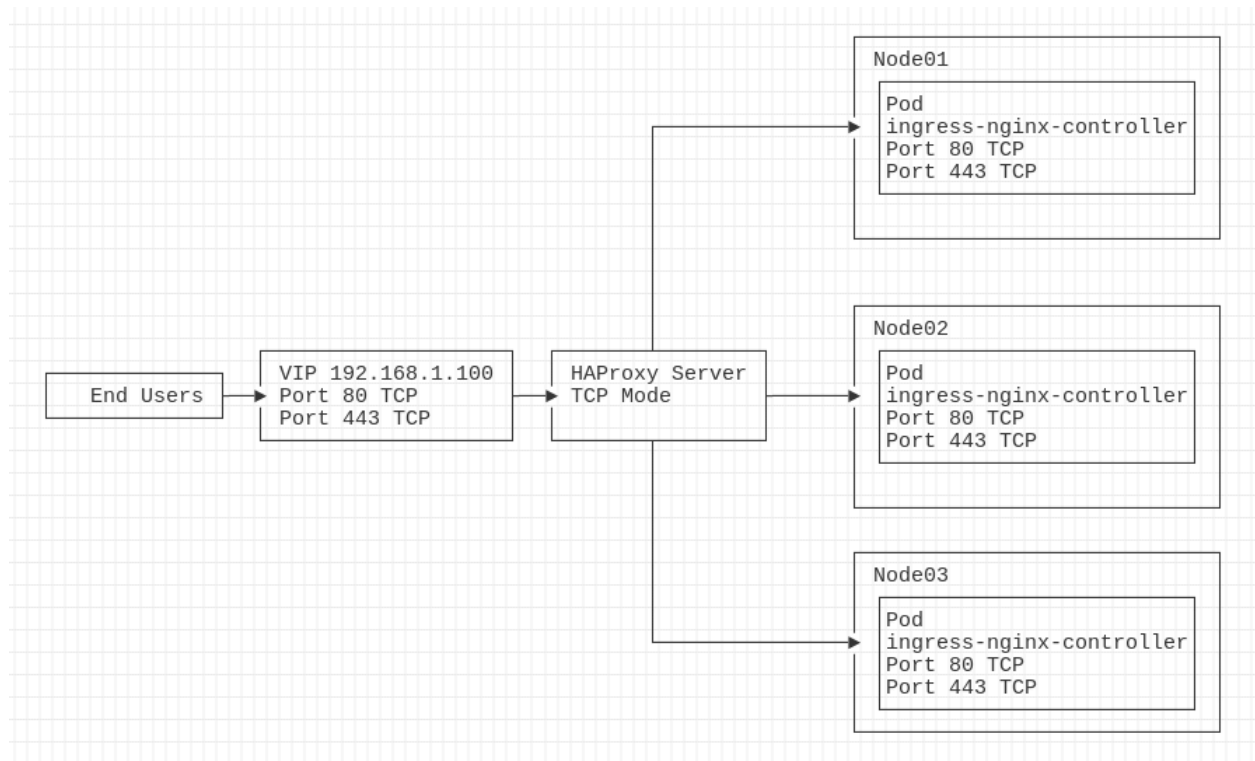
services:
  kube-api:
    audit_log:
      enabled: true
  etcd:
    backup_config:
      enabled: true           # enables recurring etcd snapshots
      interval_hours: 3      # time increment between snapshots
      retention: 72          # time in days before snapshot purge
      # Optional S3
      s3backupconfig:
        access_key: "ABCDE...."
        secret_key: "12345679...."
        bucket_name: "etcd"
        folder: "examplecluster"
        endpoint: "s3.us-west-1.amazonaws.com"
        region: "us-west-1"
```









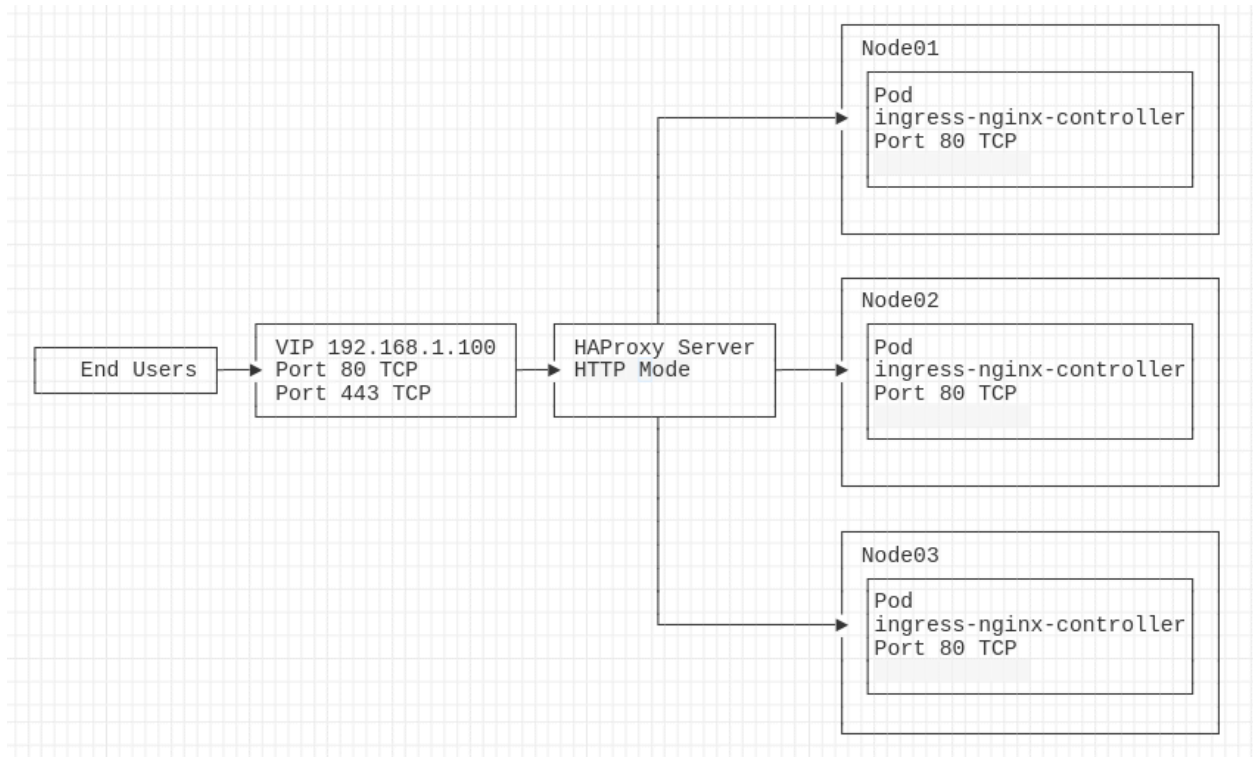


```
frontend prometheus
    bind *:8404
    option http-use-htx
    http-request use-service prometheus-exporter if { path /metrics }
    stats enable
    stats uri /stats
    stats refresh 10s

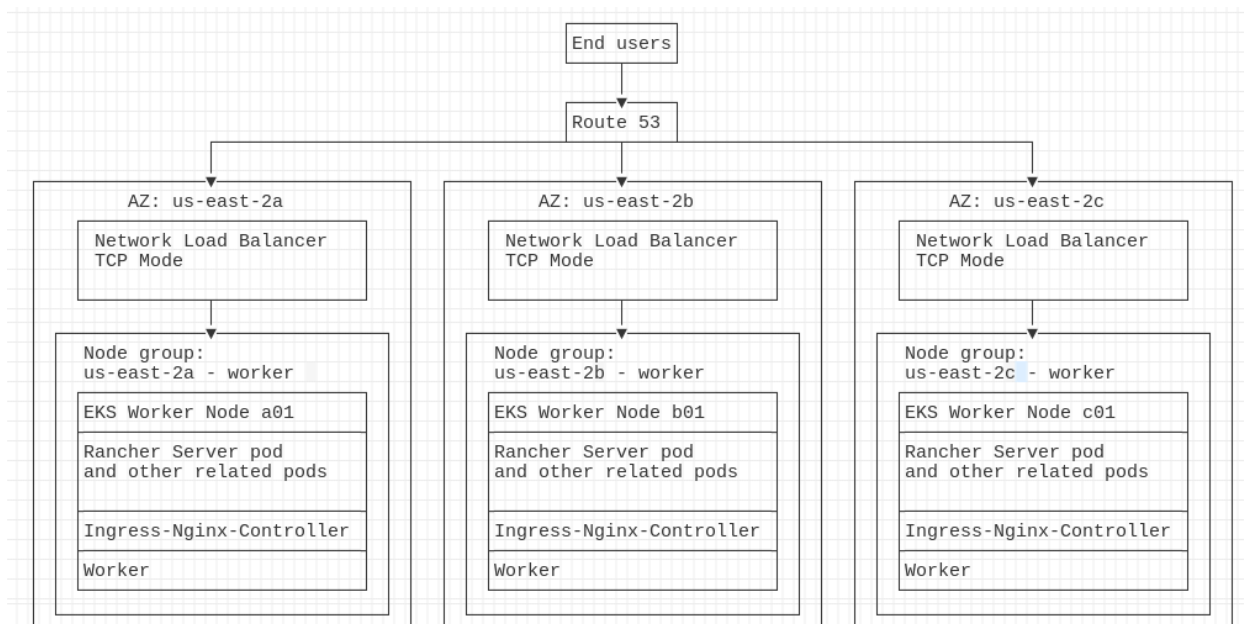
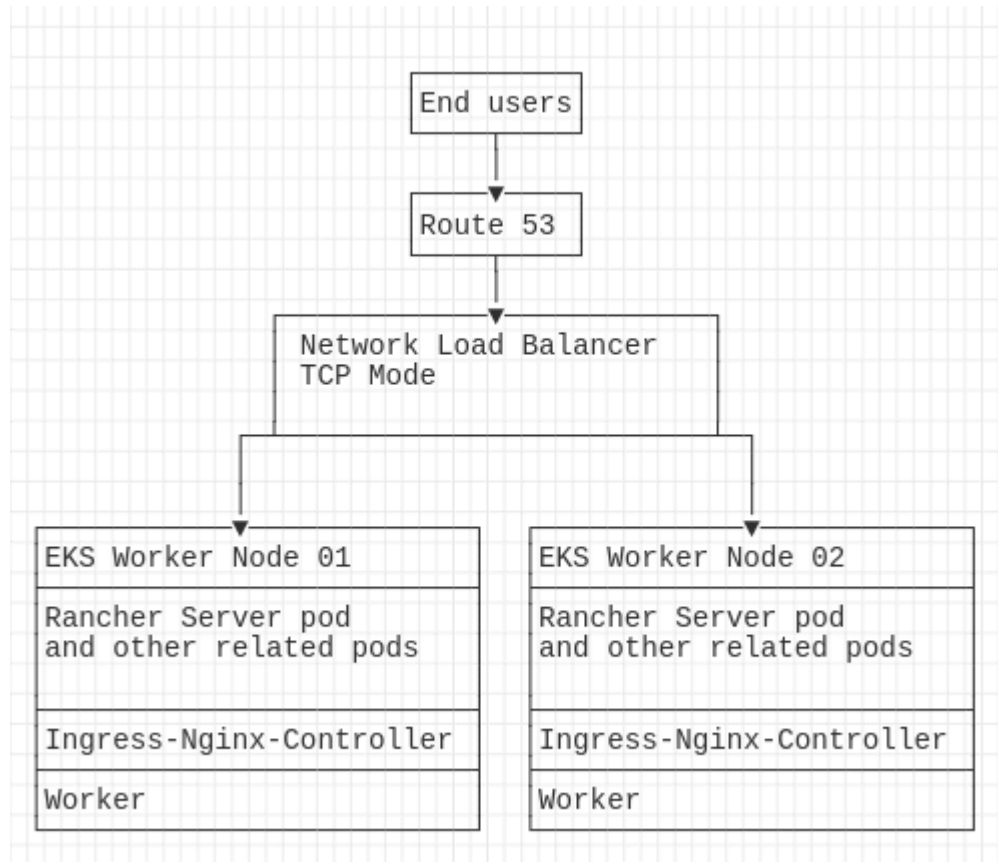
listen stats
    bind :9000
    mode http
    stats enable
    stats hide-version
    stats realm Haproxy\ Statistics
    stats uri /
    stats auth admin:Passw0rd

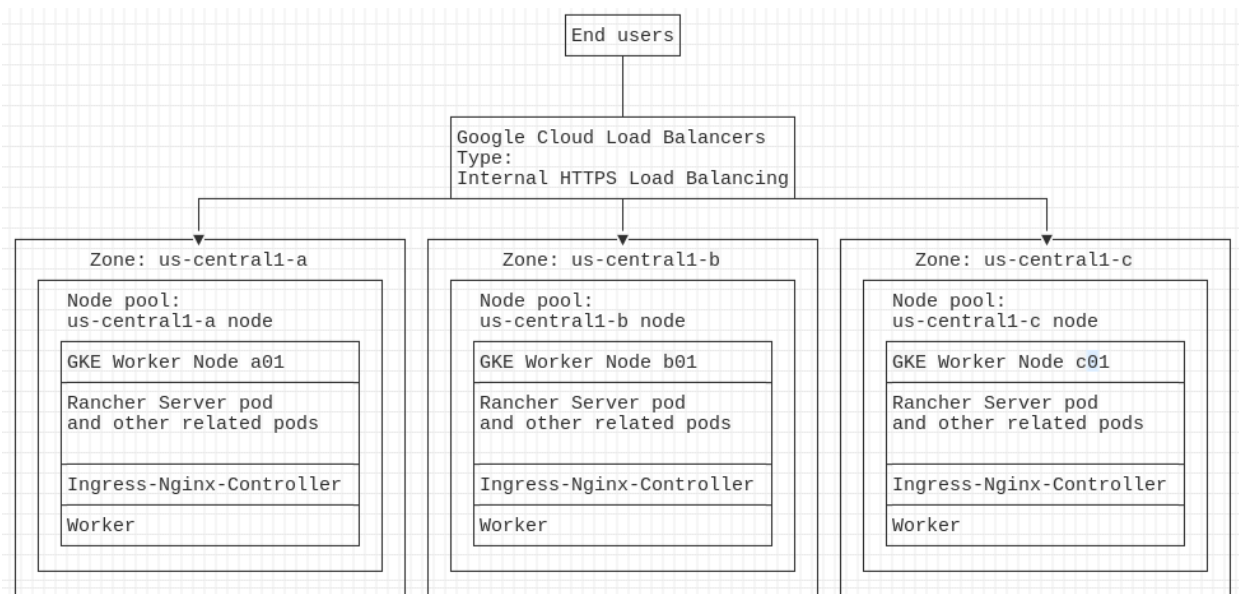
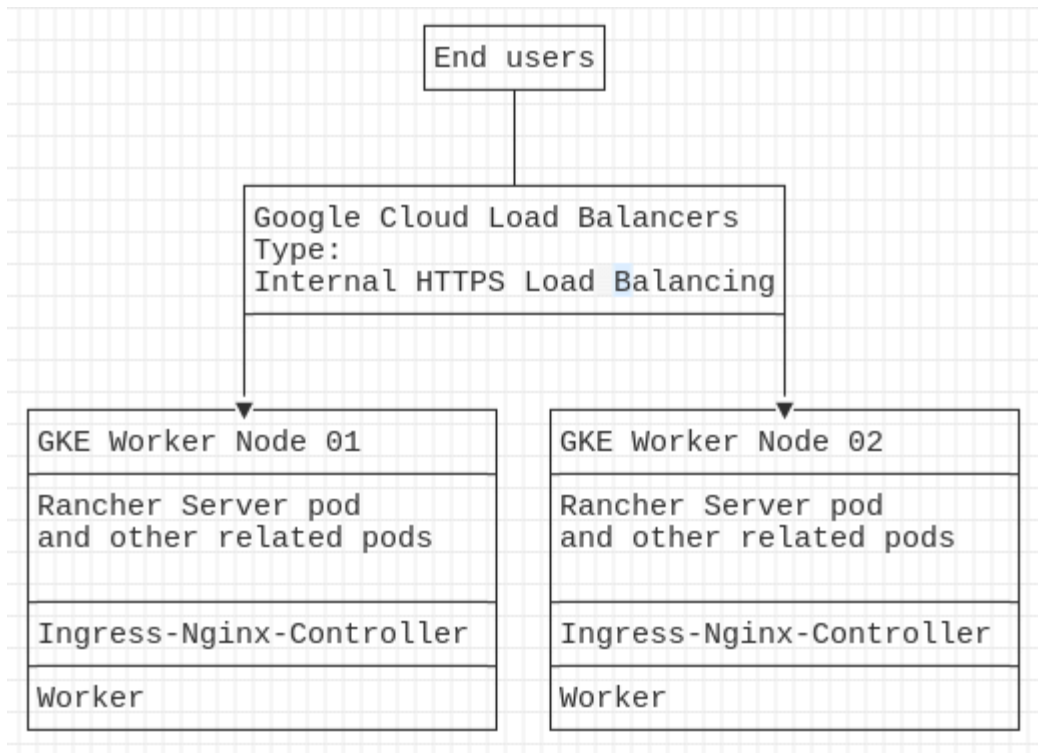
listen http
    bind *:80
    mode tcp
    balance leastconn
    stick match src
    stick-table type ip size 200k expire 30m
    server node01 192.168.1.101:80 check
    server node02 192.168.1.102:80 check
    server node03 192.168.1.103:80 check

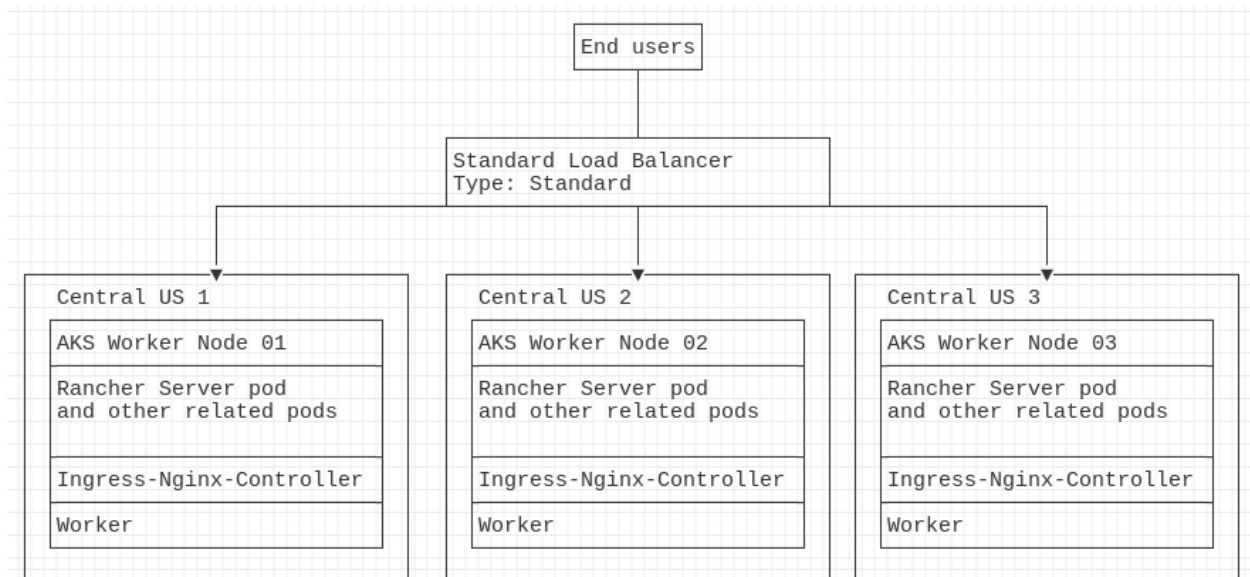
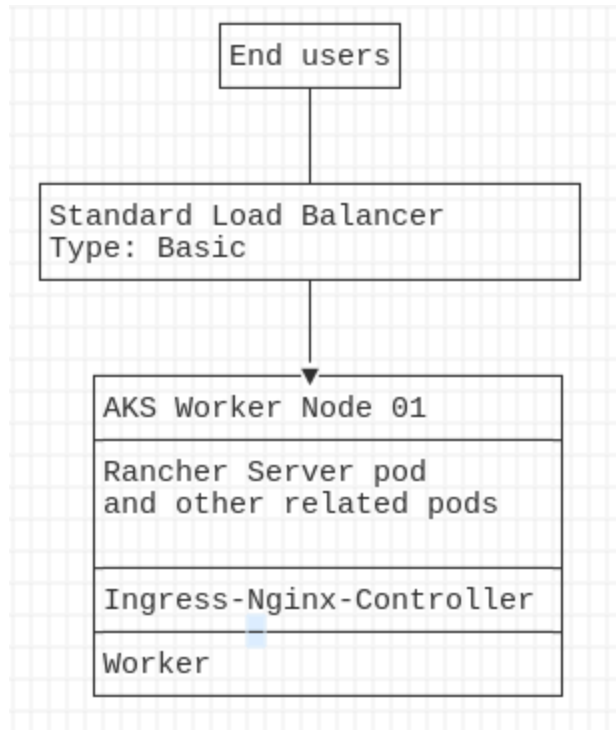
listen https
    bind *:443
    mode tcp
    balance leastconn
    stick match src
    stick-table type ip size 200k expire 30m
    server node01 192.168.1.101:443 check ssl verify none
    server node02 192.168.1.102:443 check ssl verify none
    server node03 192.168.1.103:443 check ssl verify none
```



Chapter 5: Deploying Rancher on a Hosted Kubernetes Cluster







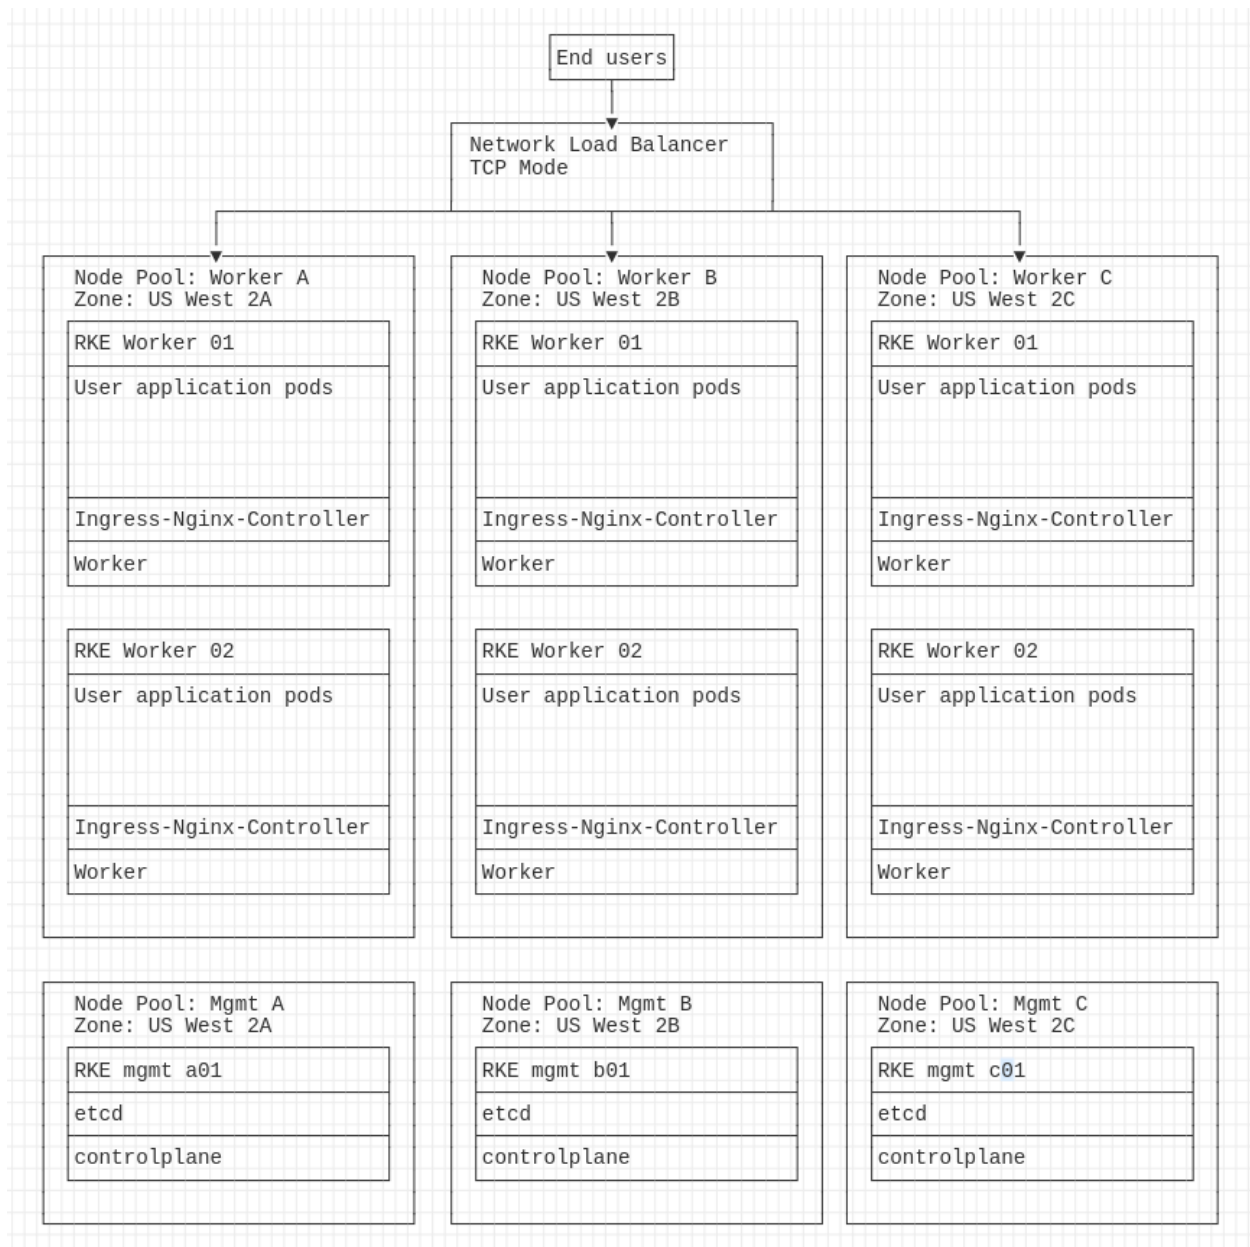
Chapter 6: Creating an RKE Cluster Using Rancher

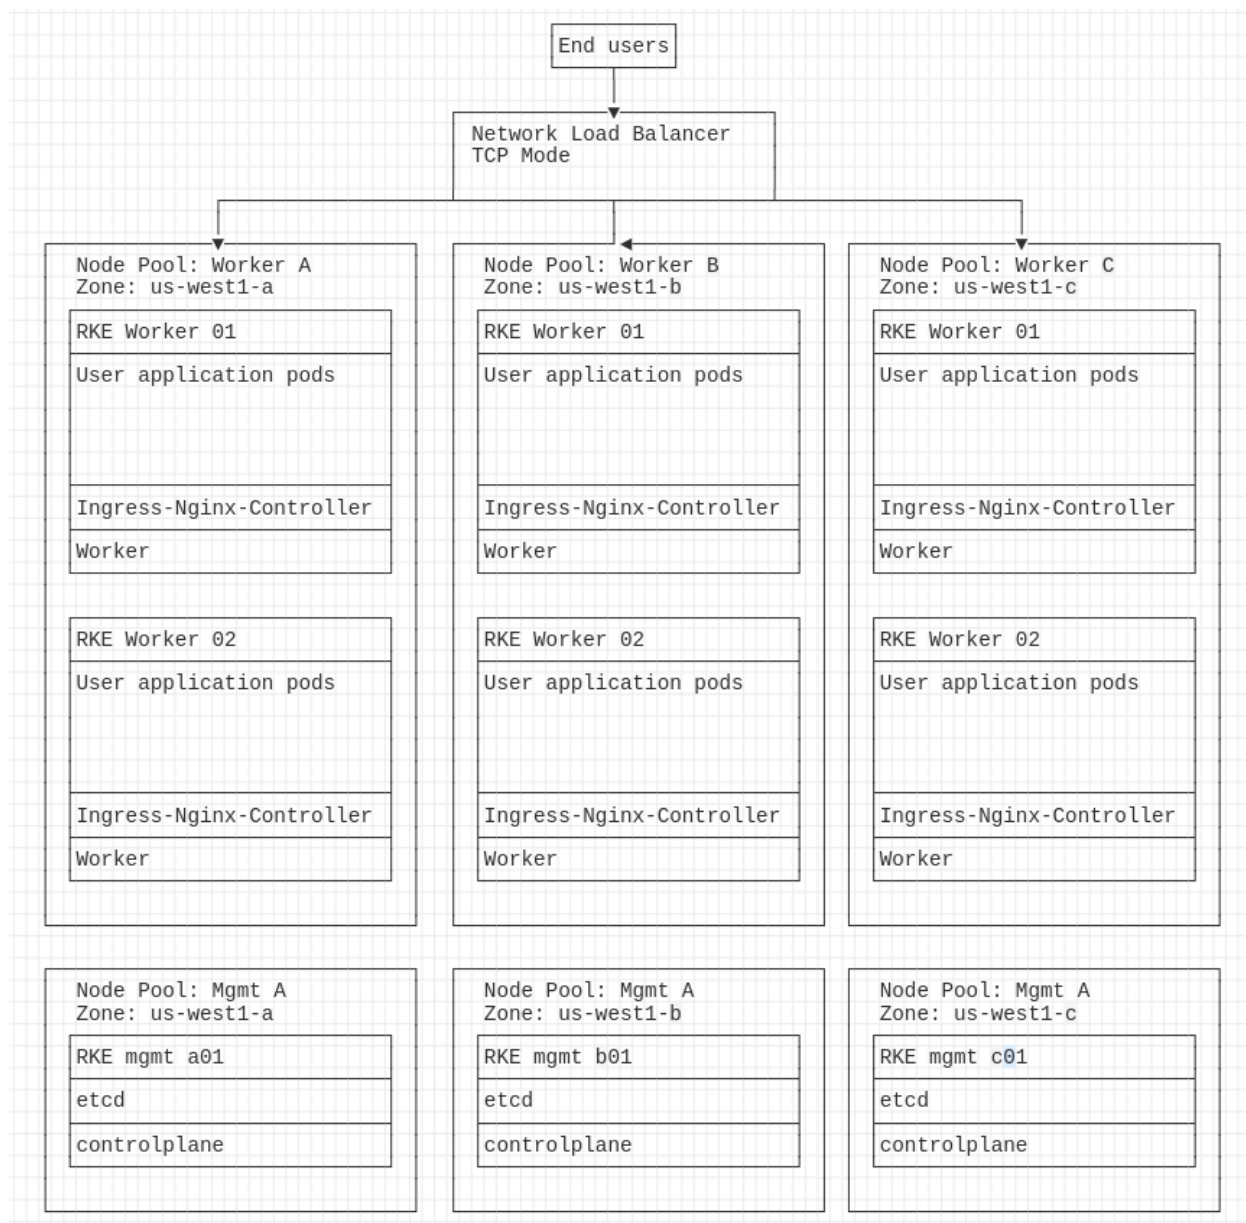
- 2 Run this command on one or more existing machines already running a supported version of Docker.

```
sudo docker run -d --privileged --restart=unless-stopped --net=host -v /etc/kubernetes:/etc/kubernetes -v /var/run:/var/run rancher/rancher-agent:v2.4.4 --server https://example.rancher.com --token qkg67qfz2lqx8p4xg2gr2czt5ft52xqmmqtjfhwldtq52pb6s645lj --ca-checksum 2bf37382990ec0c3f3539849577a122126040dcca01556231b0eea813817d --worker
```

- 2 Run this command on one or more existing machines already running a supported version of Docker.

```
sudo docker run -d --privileged --restart=unless-stopped --net=host -v /etc/kubernetes:/etc/kubernetes -v /var/run:/var/run rancher/rancher-agent:v2.4.4 --server https://example.rancher.com --token qkg67qfz2lqx8p4xg2gr2czt5ft52xqmmqtjfhwldtq52pb6s645lj --ca-checksum 2bf37382990ec0c3f3539849577a122126040dcca01556231b0eea813817d --node-name example-hostname --address 1.1.1.1 --internal-address 10.0.0.1 --etcd --controlplane --worker
```





Cloud Credential: Create Amazon

Name *

A unique name

Description

Any text you want that better describes this resource

Access Key

Your AWS Access Key

SecretKey

Your AWS Secret Key

Default Region

us-west-2



The default region to use when creating clusters. Also contacted to verify that this credential works.

Cancel

Create

Add Node Template



Amazon EC2



Azure



DigitalOcean



Rackspace



vSphere

AMAZON EC2 OPTIONS

1. Account Access

Choose the region and API Key that will be used to launch EC2 Instances

Region

us-west-2



Cloud Credentials

Rancher Example



Add New

Next: Authenticate & configure nodes

Cancel

▼ 4. Instance

Customize the EC2 Instance that will be created.

Instance Type

t2.medium

Spot Instance

☐ Request spot instance

Root Disk Size

16

GB

Encryption

☐ Encrypt EBS Volume

HTTP Endpoint

☒ Enabled ☐ Disabled

Enables or disables the HTTP metadata endpoint on your instances

HTTP Tokens

☐ Required ☒ Optional

Use HTTP Tokens for Instance Metadata Requests

AMI

An Ubuntu AMI

IAM Instance Profile Name

my-k8s-profile

SSH User

ubuntu

Private IP

☐ Use only private IP address

+ Add AWS Tag

Name *

Description

Name

Description

▼ Labels

Labels are key/value pairs that can be used to annotate containers and make scheduling decisions.

+ Add Label

▶ Taints

Taints mark that the node should not accept any pods that do not tolerate the taints.

▶ Engine Options

Customize the configuration of the Docker daemon

Create

Cancel

Cluster: Create

Create a cluster in a hosted Kubernetes provider



Amazon EKS



Azure AKS



Google GKE

Provision new nodes and create a cluster using RKE

RKE1 ☒ RKE2/K3s



Amazon EC2



Azure



DigitalOcean



RackSpace



VMware vSphere

Use existing nodes and create a cluster using RKE



Custom

Add Cluster - Custom

Cluster Name *


[Add a Description](#)

Example Name

- ▶ **Member Roles**
Control who has access to the cluster and what permission they have to change it.

- ▶ **Labels & Annotations**
Configure labels and annotations for the cluster. None

Cluster Options

[Edit as YAML](#) 

☐ Use an existing RKE Template and revision

[Expand All](#)

- ▶ **Kubernetes Options**
Customize the kubernetes cluster options

- ▶ **Private Registry**
Configure a default private registry for provisioning RKE cluster components. When enabled, all system images required for RKE cluster provisioning and system add-ons startup will be pulled from this registry.

- ▶ **Advanced Options**
Customize advanced cluster options

- ▶ **Authorized Endpoint**
Enabling the authorized cluster endpoint allows direct communication with the cluster, bypassing the API proxy. Authorized endpoints can be retrieved by generating a kubeconfig for the cluster.

[Next](#)

[Cancel](#)

Cluster Options

Customize Node Run Command

Editing node options will update the command you will run on your existing machines

1 Node Options

Choose what roles the node will have in the cluster.

Node Role

☐

etcd

☐

Control Plane

☒

Worker



The cluster needs to have at least one node with each role in order for A1 Rancher Prod to finish provisioning.

Show advanced options

2

Run this command on one or more existing machines already running a supported version of Docker.

```
sudo docker run -d --privileged --restart=unless-stopped --net=host -v /etc/kubernetes:/etc/kubernetes -v /var/run:/var/run rancher/rancher-agent:v2.6.2 --server https://example.rancher.com --token kfn7m6qp7gkjh2l7m7tqjfj79b5dcfs6gbc62zwf2p7s6m9gwwmzsq --worker
```



Name Prefix ▾ Count ▾ Template ▾ Auto Replace ▾ Drain Before Delete etcd Control Plane Worker Taints

1

I ▾ +

0

minutes

☐☐☐☒

Taints

−

Number of nodes required:

⊗ 1, 3, or 5 ⊗ 1 or more ✔ 1 or more

+ Add Node Pool

Chapter 7: Deploying a Hosted Cluster with Rancher

Cloud Credential: Create Amazon

Name * A unique name	Description Any text you want that better describes this resource
--------------------------------	---

Access Key Your AWS Access Key
--

SecretKey Your AWS Secret Key

Default Region us-west-2	▼
------------------------------------	---

The default region to use when creating clusters. Also contacted to verify that this credential works.

Cloud Credential: Create Google

Name * A unique name	Description Any text you want that better describes this resource
--------------------------------	---

Service Account Service Account private key JSON file

Create a **Service Account** with a JSON private key and provide the JSON here. These IAM roles are required:

- Compute Engine: Compute Viewer (roles/compute.viewer)
- Project: Viewer (roles/viewer)
- Kubernetes Engine: Kubernetes Engine Admin (roles/container.admin)
- Service Accounts: Service Account User (roles/iam.serviceAccountUser)

More info on roles can be found [here](#).

Cloud Credential: Create Azure

Name * A unique name	Description Any text you want that better describes this resource
Environment * AzurePublicCloud	Subscription ID *
Client ID *	Client Secret *

[Cancel](#) [Create](#)

Add Cluster - Amazon EKS

Cluster Name * [Add a Description](#)

e.g. sandbox

▶ **Member Roles**
Control who has access to the cluster and what permission they have to change it.

▶ **Labels & Annotations**
Configure labels and annotations for the cluster. [None](#)

[Expand All](#)

▼ **Account Access**
Choose the region and API Key that will be used to launch Amazon EKS

Region	Cloud Credentials	
us-west-2	AWS-RancherLabs	Add New

[Next: Configure Cluster](#) [Cancel](#)

Chapter 8: Importing an Externally Managed Cluster into Rancher

No Images...

Chapter 9: Cluster Configuration Backup and Recovery

No Images...

Chapter 10: Monitoring and Logging

The screenshot shows the Rancher dashboard interface. At the top, a dark navigation bar contains the Rancher logo and a menu with items: local, Cluster, Nodes, Storage, Projects/Namespaces, Members, and Tools. The 'Tools' menu is expanded, showing a list of options: Alerts, Catalogs, Notifiers, Logging, Monitoring (highlighted in blue), Istio, CIS Scans, Rancher Backups, and OPA Gatekeeper. Below the navigation bar, the main content area is titled 'Dashboard: local'. A light blue banner with an information icon and text reads: 'Try out the Cluster Explorer for a new way to view and manage your Kubernetes resources'. Below this, a light blue box indicates 'Provider: Imported'. The main content area is divided into two columns by a large, light gray curved line. The left column contains a warning message about legacy V1 Istio Helm v1.5 and a 'Cluster Tools' button. The right column displays five installation cards: 'Legacy V1 Istio Helm v1.5', 'Longhorn', 'Monitoring', 'OPA Gatekeeper', and 'Banzai Cloud Logging Operator'. Each card includes a description, version information, and an 'Install' button.

local Cluster Nodes Storage Projects/Namespaces Members Tools

Alerts
Catalogs
Notifiers
Logging
Monitoring
Istio
CIS Scans
Rancher Backups
OPA Gatekeeper

Dashboard: local

Try out the [Cluster Explorer](#) for a new way to view and manage your Kubernetes resources

Provider: Imported

Legacy V1 Istio Helm v1.5 has been deprecated since Rancher 2.5.0. [Learn more](#) about migrating to the latest versions. [Install](#)

Collects and filter logs using highly configurable CRDs. Powered by Banzai Cloud Logging Operator. [Install](#)

Longhorn
v1.0.0 (1.0.0-2.0.0)
Longhorn is a distributed block storage system for Kubernetes. [Install](#)

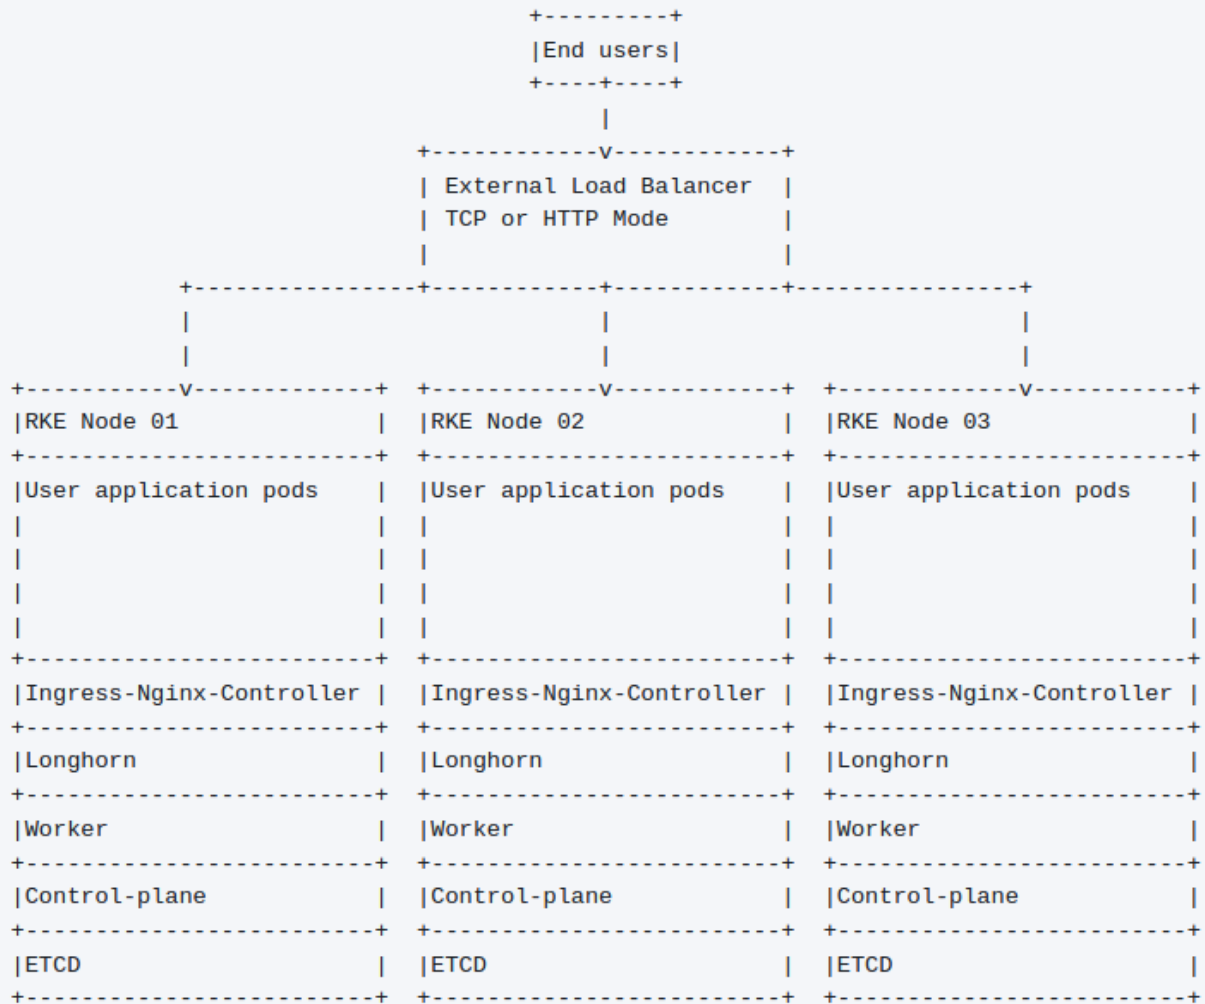
Monitoring
v0.0.0 (0.0.0-0.0.0)
Collects several related Helm charts, Grafana dashboards, and Prometheus rules combined with documentation and scripts to provide ease to operate end-to-end Kubernetes cluster monitoring with Prometheus using the Prometheus Operator. [Install](#)

OPA Gatekeeper
v0.0.0 (0.0.0-0.0.0)
Modifies Open Policy Agent's upstream gatekeeper chart that provides policy-based control for cloud native environments. [Install](#)

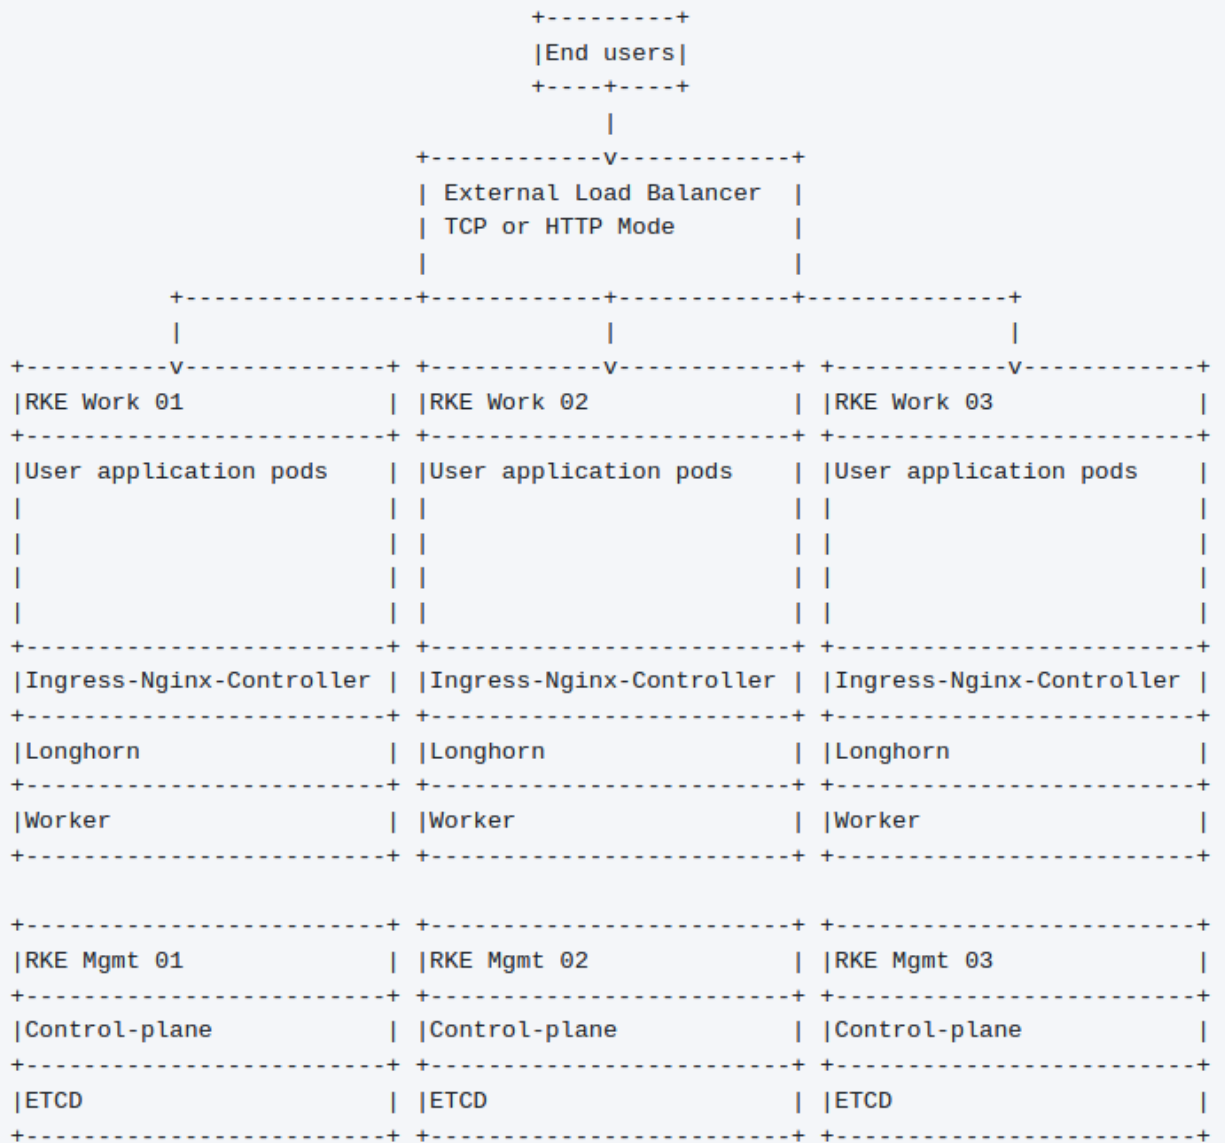
Cluster Tools

v2.4.0

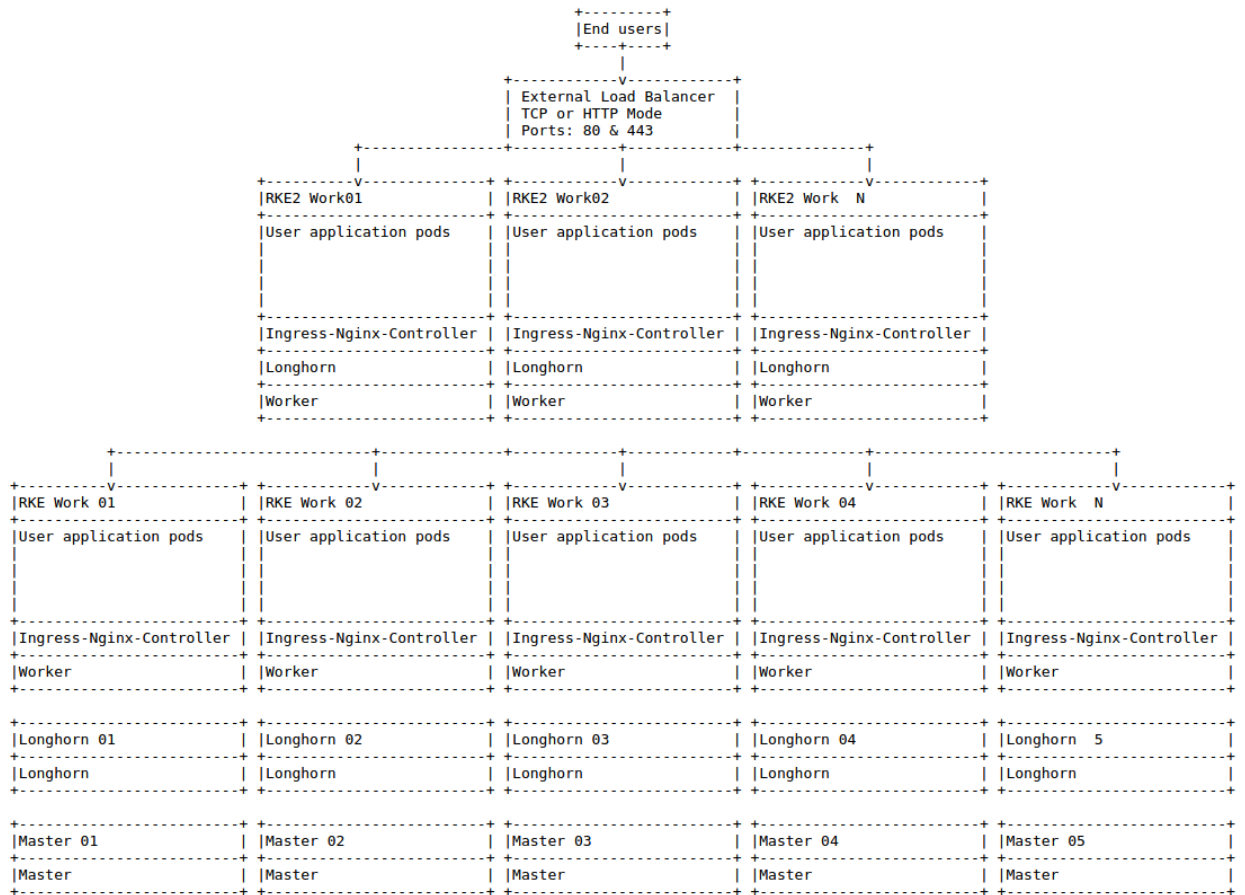
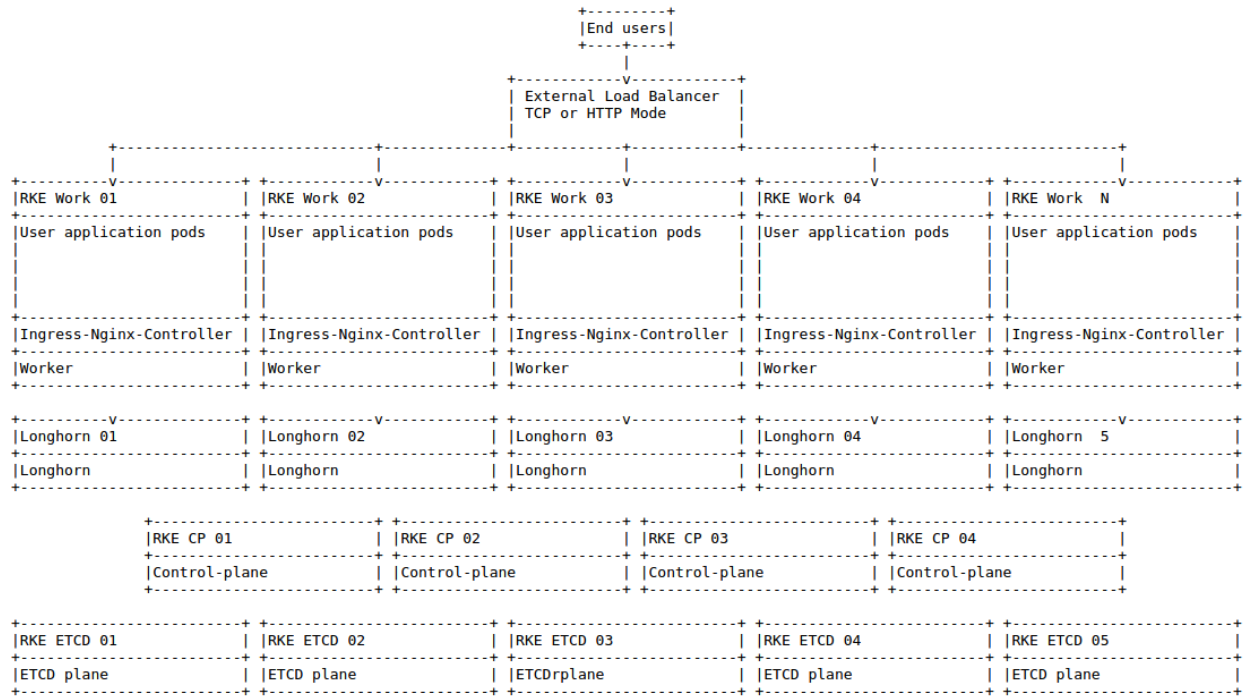
Chapter 11: Bring Storage to Kubernetes



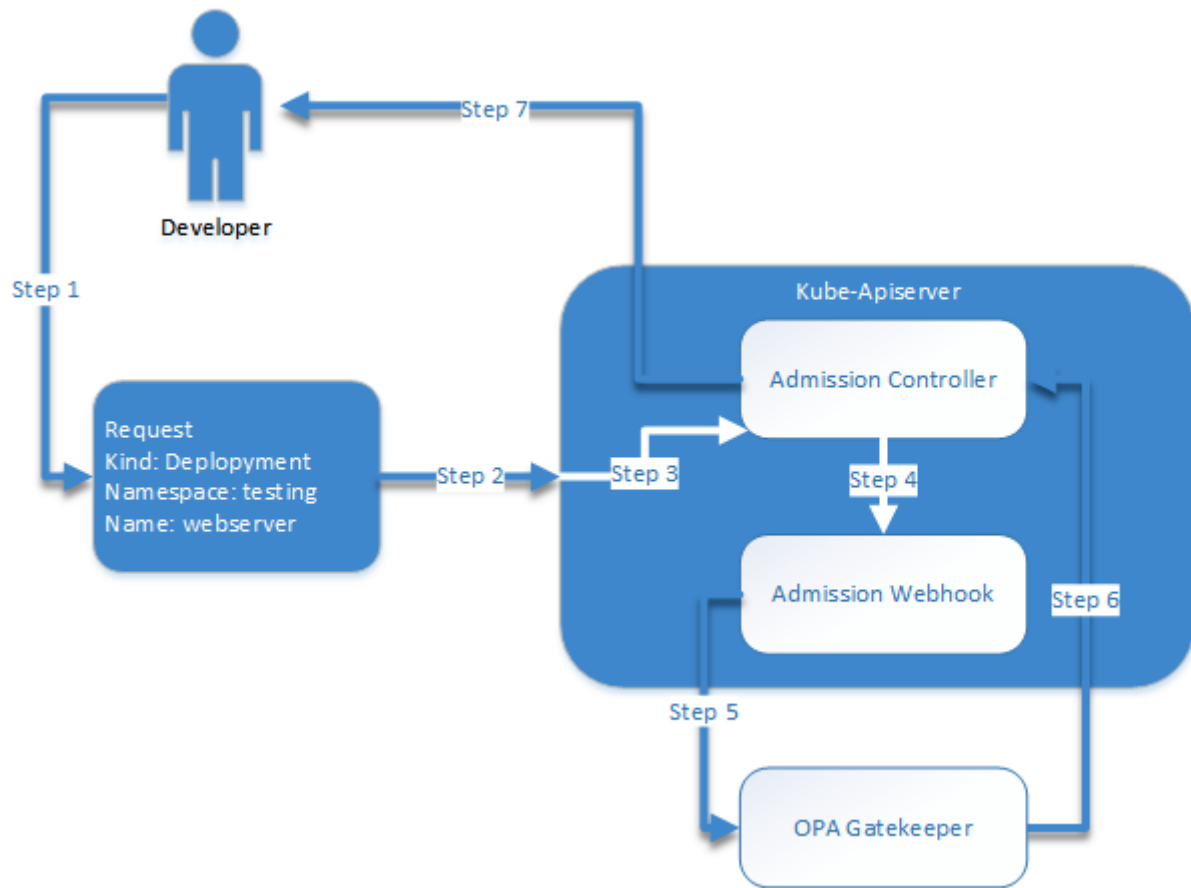








Chapter 12: Security and Compliance (OPA Gatekeeper)



er-lab

er for a new way to view and manage your Kubernetes resources.

ider: ImportedKubernetes

%

1%

Alerts

Catalogs

Notifiers

Logging

Monitoring

Istio

CIS Scans

Rancher Backups ⌘

OPA Gatekeeper ⌘

Install OPA Gatekeeper



Chart
OPA Gatekeeper

Version
3.3.001

Description
Any text you want that better describes this resource

Install into Project
(None)

README

Values YAML

Helm README

Helm Deploy Options

README

Rancher OPA Gatekeeper

This chart is based off of the upstream [OPA Gatekeeper](#) chart. For more information on how to use the feature, refer to our [docs](#). The chart installs the following components:

- **OPA Gatekeeper Controller-Manager** - OPA Gatekeeper is a policy engine for providing policy based governance for Kubernetes clusters. The controller installs as a [validating admission controller webhook](#) on the cluster and intercepts all admission requests that create, update or delete a resource in the cluster.
- **Audit** - A periodic audit of the cluster resources against the enforced policies. Any existing resource that violates a policy will be recorded as violations.
- **Constraint Template** - A template is a CRD (`ConstraintTemplate`) that defines the schema and Rego logic of a policy to be applied to the cluster by Gatekeeper's admission controller webhook. This chart installs a few default `ConstraintTemplate` custom resources.
- **Constraint** - A constraint is a custom resource that defines the scope of resources which a specific constraint template should apply to. The complete policy is defined by a combination of `ConstraintTemplates` (i.e. what the policy is) and `Constraints` (i.e. what resource to apply the policy to).

For more information on how to configure the Helm chart, refer to the Helm README.

Cancel

Install

Apps & Marketplace

Cluster Manager

Charts

Installed Apps7

Chart Repositories2

Recent Operations0

Deploy Chart

All

Rancher

Partners

All Categories

Filter

Alerting Drivers

The manager for third-party webhook receivers used in Prometheus Alertmanager

CIS Benchmark

The cis-operator enables running CIS benchmark security scans on a Kubernetes cluster

External IP Webhook

Deploy the external-ip-webhook to mitigate k8s CVE-2020-8554

Istio

A basic Istio setup that installs with the istioctl. Refer to <https://istio.io/latest/> for details.

Logging

Collects and filter logs using highly configurable CRDs. Powered by Banzai Cloud Logging Operator.

Longhorn

Longhorn is a distributed block storage system for Kubernetes.

Monitoring

Collects several related Helm charts, Grafana dashboards, and Prometheus rules combined with documentation and script...

OPA Gatekeeper

Modifies Open Policy Agent's upstream gatekeeper chart that provides policy-based control for cloud native...

Rancher Backups

Provides ability to back up and restore the Rancher application running on any

vSphere CPI

vSphere Cloud Provider Interface (CPI)

vSphere CSI

vSphere Cloud Storage Interface (CSI)

Ambassador Edge Stack

A Helm chart for Datawire Ambassador

Starred

Cluster

Workload

Apps & Marketplace

Charts

Installed Apps13

Repositories2

Recent Operations0

Service Discovery

Storage

RBAC

Charts

All

OPA Gatekeeper

Modifies Open Policy Agent's upstream gatekeeper chart that provides policy-based control for cloud native environments



Charts: OPA Gatekeeper (100.0.1+up3.6.0)

[Install](#)

Modifies Open Policy Agent's upstream gatekeeper chart that provides policy-based control for cloud native environments

This chart is based off of the upstream [OPA Gatekeeper](#) chart.
For more information on how to use the feature, refer to our [docs](#).
The chart installs the following components:

- OPA Gatekeeper Controller-Manager - OPA Gatekeeper is a policy engine for providing policy based governance for Kubernetes clusters. The controller installs as a [validating admission controller webhook](#) on the cluster and intercepts all admission requests that create, update or delete a resource in the cluster.
- [Audit](#) - A periodic audit of the cluster resources against the enforced policies. Any existing resource that violates a policy will be recorded as violations.
- [Constraint Template](#) - A template is a CRD ([ConstraintTemplate](#)) that defines the schema and Rego logic of a policy to be applied to the cluster by Gatekeeper's admission controller webhook. This chart installs a few default [ConstraintTemplate](#) custom resources.
- [Constraint](#) - A constraint is a custom resource that defines the scope of resources which a specific constraint template should apply to. The complete policy is defined by a combination of [ConstraintTemplates](#) (i.e. what the policy is) and [Constraints](#) (i.e. what resource to apply the policy to).

Chart Versions

100.0.1+up3.6.0

Wed, Jan 5 2022

Application Version

v3.6.0

Home

<https://github.com/open-policy-agent/gatekeeper>

Related

<https://github.com/open-policy-agent/gatekeeper.git>



OPA Gatekeeper
100.0.1+up3.6.0

Install: Step 1
Set App metadata

Metadata

This process will help create the chart. Start by setting some basic information used by A1 Rancher Prod to manage the App.

Install into Project
System

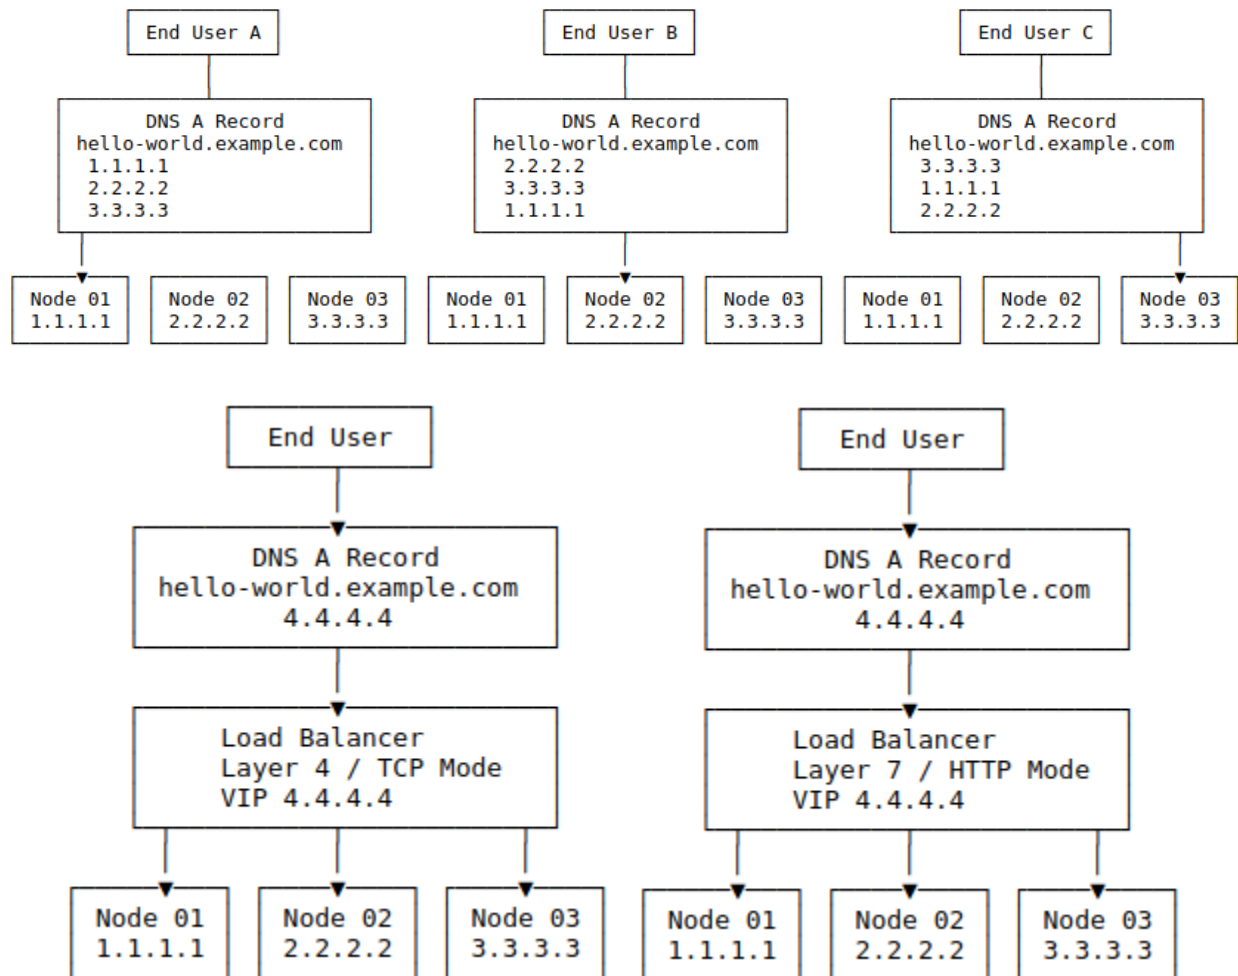


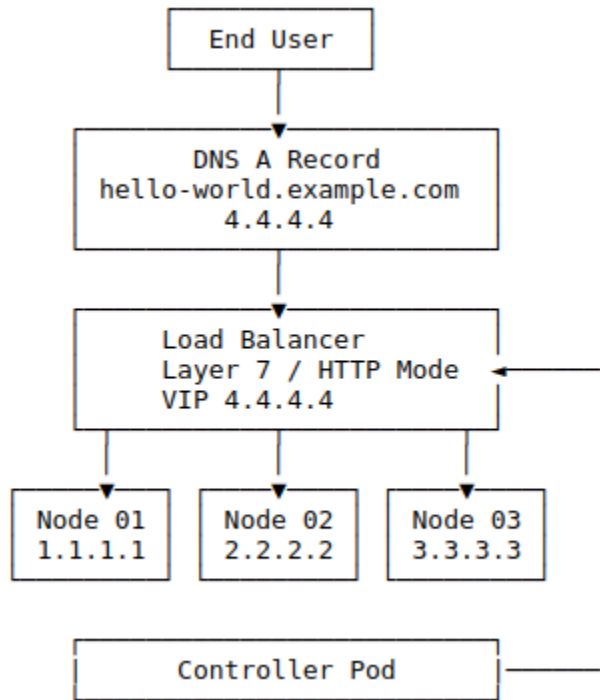
Chapter 13: Scaling in Kubernetes

Reason	Age	From	Message
-----	----	----	-----
SuccessfulRescale	11m	horizontal-pod-autoscaler	New size: 8; reas
SuccessfulRescale	11m	horizontal-pod-autoscaler	New size: 10; rea
SuccessfulRescale	10m (x2 over 11m)	horizontal-pod-autoscaler	New size: 4; reas
SuccessfulRescale	9m53s	horizontal-pod-autoscaler	New size: 8; reas
SuccessfulRescale	9m37s	horizontal-pod-autoscaler	New size: 10; rea
FailedGetResourceMetric	4m19s	horizontal-pod-autoscaler	failed to get cpu
FailedComputeMetricsReplicas	4m18s	horizontal-pod-autoscaler	invalid metrics (
ics returned from resource metrics API			
FailedGetResourceMetric	3m18s (x6 over 10m)	horizontal-pod-autoscaler	failed to get cpu
FailedComputeMetricsReplicas	3m18s (x6 over 10m)	horizontal-pod-autoscaler	invalid metrics (
SuccessfulRescale	3m3s	horizontal-pod-autoscaler	New size: 1; reas

```
Status:
Conditions:
  Last Transition Time:  2022-02-20T11:08:45Z
  Status: True
  Type: RecommendationProvided
Recommendation:
  Container Recommendations:
    Container Name: hello-world
    Lower Bound:
      Cpu: 12m
      Memory: 247917233
    Target:
      Cpu: 63m
      Memory: 380258472
    Uncapped Target:
      Cpu: 63m
      Memory: 380258472
    Upper Bound:
      Cpu: 137m
      Memory: 561393834
    Container Name: hello-world
    Lower Bound:
      Cpu: 12m
      Memory: 131072k
    Target:
      Cpu: 12m
      Memory: 131072k
    Uncapped Target:
      Cpu: 12m
      Memory: 131072k
    Upper Bound:
      Cpu: 16m
      Memory: 131072k
Events: <none>
```


Chapter 14: Load Balancer Configuration and SSL Certificates





Configuration: Basic ▾

Name	<input type="text" value="Rancher_80"/>
Description	<input type="text"/>
Health Monitors	<div><div>Active</div><div><div>/Common</div><div>http</div></div></div> <div><< >></div> <div><div>Available</div><div><div>/Common</div><div>gateway_icmp</div><div>http2</div><div>http2_head_f5</div><div>http_head_f5</div></div></div>

Resources

Load Balancing Method	<input type="text" value="Round Robin"/> ▾																				
Priority Group Activation	<input type="text" value="Disabled"/> ▾																				
New Members	<div><div><input checked="" type="radio"/> New Node <input type="radio"/> New FQDN Node</div><div>Node Name: <input type="text" value="node03"/> (Optional)</div><div>Address: <input type="text" value="192.168.1.103"/></div><div>Service Port: <input type="text" value="80"/> <input type="text" value="HTTP"/> ▾</div><div><input type="button" value="Add"/></div><table><thead><tr><th>Node Name</th><th>Address/FQDN</th><th>Service Port</th><th>Auto Populate</th><th>Priority</th></tr></thead><tbody><tr><td>node01</td><td>192.168.1.101</td><td>80</td><td></td><td>0</td></tr><tr><td>node02</td><td>192.168.1.102</td><td>80</td><td></td><td>0</td></tr><tr><td>node03</td><td>192.168.1.103</td><td>80</td><td></td><td>0</td></tr></tbody></table><div><input type="button" value="Edit"/> <input type="button" value="Delete"/></div></div>	Node Name	Address/FQDN	Service Port	Auto Populate	Priority	node01	192.168.1.101	80		0	node02	192.168.1.102	80		0	node03	192.168.1.103	80		0
Node Name	Address/FQDN	Service Port	Auto Populate	Priority																	
node01	192.168.1.101	80		0																	
node02	192.168.1.102	80		0																	
node03	192.168.1.103	80		0																	

General Properties

Name	Rancher_80
Partition / Path	Common
Description	<input type="text"/>
Type	Performance (Layer 4) ▼
Source Address	<input checked="" type="radio"/> Host <input type="radio"/> Address List <input type="text" value="0.0.0.0/0"/>
Destination Address/Mask	<input checked="" type="radio"/> Host <input type="radio"/> Address List <input type="text" value="192.168.1.99"/>
Service Port	<input checked="" type="radio"/> Port <input type="radio"/> Port List <input type="text" value="80"/> HTTP ▼
Notify Status to Virtual Address	<input checked="" type="checkbox"/>
Availability	<input checked="" type="checkbox"/> Unknown (Enabled) - The children pool member(s) either don't have service check
Syncookie Status	Inactive
State	Enabled ▼

Configuration: Basic ▼

DoH Profile Type	None
Protocol	TCP ▼
Protocol Profile (Client)	fastL4 ▼
HTTP Profile (Client)	None ▼
HTTP Profile (Server)	(Use Client Profile) ▼
HTTP Proxy Connect Profile	None ▼
VLAN and Tunnel Traffic	All VLANs and Tunnels ▼
Source Address Translation	None ▼

Local Traffic » Virtual Servers : Virtual Server List » Rancher_80

Properties

Load Balancing

Default Pool	<input type="text" value="Rancher_80"/>
Default Persistence Profile	<input type="text" value="None"/>
Fallback Persistence Profile	<input type="text" value="None"/>

Debian/Ubuntu HAProxy packages

The Debian HAProxy packaging team provides various versions of [HAProxy](#) packages for use on different Debian or Ubuntu systems. The following wizard helps you to find the package suitable for your system.

<< I am running and I want to install HAProxy
 version .

Instructions for latest release

You need to enable a [dedicated PPA](#) with the following command:

```
# apt-get install --no-install-recommends software-properties-common
# add-apt-repository ppa:vbernat/haproxy-2.5
```

Then, use the following command:

```
# apt-get install haproxy=2.5.*
```

You will get the *latest* release of HAProxy 2.5 (and stick to this branch).

```
frontend http_frontend
  mode tcp
  bind :80
  default_backend http_backend

backend http_backend
  mode tcp
  balance leastconn
  server node01 192.168.1.101:80 check
  server node02 192.168.1.102:80 check
  server node03 192.168.1.103:80 check

frontend https_frontend
  mode tcp
  bind :80
  default_backend https_backend

backend https_backend
  mode tcp
  balance leastconn
  server node01 192.168.1.101:443 check
  server node02 192.168.1.102:443 check
  server node03 192.168.1.103:443 check
```

```
frontend frontend:
  bind *:80
  bind *:443 ssl crt /etc/haproxy/certs/star.example.com.pem
  http-request redirect scheme https unless { ssl_fc }
  mode http

acl host_rke-cluster-npd hdr(host) -i dev.example.com
acl host_rke-cluster-npd hdr(host) -i gas.example.com
use_backend rke-cluster-npd-https if host_a1-rke-cluster-npd

acl host_rke-cluster-prd hdr(host) -i example.com
acl host_rke-cluster-prd hdr(host) -i www.example.com
use_backend rke-cluster-prd-https if host_a1-rke-cluster-prd
```

```
backend rke-cluster-npd-https
  mode http
  option httpchk HEAD /healthz HTTP/1.0
  server node01 192.168.1.101:443 check weight 1 maxconn 1024 ssl verify none
  server node02 192.168.1.102:443 check weight 1 maxconn 1024 ssl verify none
  server node03 192.168.1.103:443 check weight 1 maxconn 1024 ssl verify none

backend rke-cluster-prd-https
  mode http
  option httpchk HEAD /healthz HTTP/1.0
  server node04 192.168.1.104:443 check weight 1 maxconn 1024 ssl verify none
  server node05 192.168.1.105:443 check weight 1 maxconn 1024 ssl verify none
  server node06 192.168.1.106:443 check weight 1 maxconn 1024 ssl verify none
```

```
apiVersion: v1
kind: ConfigMap
metadata:
  namespace: metallb-system
  name: config
data:
  config: |
    address-pools:
    - name: default
      protocol: layer2
      addresses:
      - 192.168.2.2-192.168.2.125
```

```
apiVersion: v1
kind: ConfigMap
metadata:
  namespace: metallb-system
  name: config
data:
  config: |
    peers:
      - peer-address: 10.0.0.1
        peer-asn: 64501
        my-asn: 64500
    address-pools:
      - name: default
        protocol: bgp
        addresses:
          - 192.168.10.0/24
```

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: ingress-wildcard-host
spec:
  rules:
    - host: "foo.bar.com"
      http:
        paths:
          - pathType: Prefix
            path: "/bar"
            backend:
              service:
                name: service1
                port:
                  number: 80
    - host: "/*.foo.com"
      http:
        paths:
          - pathType: Prefix
            path: "/foo"
            backend:
              service:
                name: service2
                port:
                  number: 80
```



```
apiVersion: v1
kind: Secret
metadata:
  name: testsecret-tls
  namespace: default
data:
  tls.crt: base64 encoded cert
  tls.key: base64 encoded key
type: kubernetes.io/tls
```

```
tls:
- hosts:
  - https-example.foo.com
  secretName: testsecret-tls
```

Chapter 15: Rancher and Kubernetes Troubleshooting

```
etcd-a1ubrked01
+-----+-----+-----+-----+
|      ENDPOINT      | HEALTH |   TOOK   |  ERROR  |
+-----+-----+-----+-----+
| https://127.0.0.1:2379 |   true  | 6.96521ms |         |
+-----+-----+-----+-----+

etcd-a1ubrked02
+-----+-----+-----+-----+
|      ENDPOINT      | HEALTH |   TOOK   |  ERROR  |
+-----+-----+-----+-----+
| https://127.0.0.1:2379 |   true  | 70.546814ms |         |
+-----+-----+-----+-----+

etcd-a1ubrked03
+-----+-----+-----+-----+
|      ENDPOINT      | HEALTH |   TOOK   |  ERROR  |
+-----+-----+-----+-----+
| https://127.0.0.1:2379 |   true  | 14.07386ms |         |
+-----+-----+-----+-----+

etcd-a1ubrked04
+-----+-----+-----+-----+
|      ENDPOINT      | HEALTH |   TOOK   |  ERROR  |
+-----+-----+-----+-----+
| https://127.0.0.1:2379 |   true  | 8.447933ms |         |
+-----+-----+-----+-----+

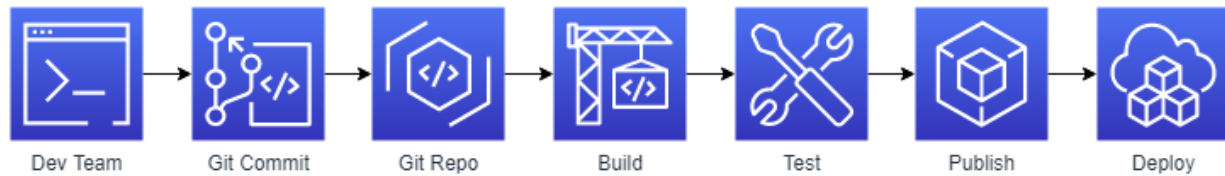
etcd-a1ubrked05
+-----+-----+-----+-----+
|      ENDPOINT      | HEALTH |   TOOK   |  ERROR  |
+-----+-----+-----+-----+
| https://127.0.0.1:2379 |   true  | 7.038ms |         |
+-----+-----+-----+-----+
```

```
apiVersion: v1
kind: Namespace
metadata:
  annotations:
    cattle.io/status: '{"Conditions":[{"Type":"InitialRolesPopulate
      resource quota","LastUpdateTime":""}, {"Type":"ResourceQuotaIn
      field.cattle.io/creatorId: u-zbw65ndtej
      field.cattle.io/projectId: c-x8rzf:p-pr7zr
      field.cattle.io/resourceQuota: "null"
    creationTimestamp: "2022-03-19T14:00:17Z"
    deletionTimestamp: "2022-03-19T14:02:05Z"
    labels:
      cattle.io/creator: norman
      field.cattle.io/projectId: p-pr7zr
    name: cattle-logging
    resourceVersion: "5872"
    selfLink: /api/v1/namespaces/t
    uid: 6a975c42-0396-11e9-bd3b-aaaaaaaaaaa4a
  spec:
    finalizers:
      - kubernetes
  status:
    phase: Terminating
```

```
apiVersion: v1
kind: Namespace
metadata:
  creationTimestamp: 2022-03-19T18:48:30Z
  deletionTimestamp: 2022-03-19T18:59:36Z
  name: test-namespace
  resourceVersion: "1385077"
  selfLink: /api/v1/namespaces/test-namespace
  uid: b50c9ea4-ec2b-11e8-a0be-fa163eeb47a5
spec:
  finalizers:
    - kubernetes
status:
  phase: Terminating
```

```
{
  "apiVersion": "v1",
  "kind": "Namespace",
  "metadata": {
    "annotations": {
      "cattle.io/status": "{\"Conditions\": [{\"Type\": \"InitialRo\",
      \"field.cattle.io/creatorId\": \"user-sw4mg\",
      \"field.cattle.io/projectId\": \"c-gkz6s:p-48tst\",
      \"kubectrl.kubernetes.io/last-applied-configuration\": \"{\\\"api\",
      \"lifecycle.cattle.io/create.namespace-auth\": \"true\"
    },
    \"creationTimestamp\": \"2022-03-01T17:17:20Z\",
    \"deletionTimestamp\": \"2022-03-18T12:30:14Z\",
    \"labels\": {
      \"cattle.io/creator\": \"norman\",
      \"field.cattle.io/projectId\": \"p-48tst\"
    },
    \"name\": \"longhorn-system\",
    \"resourceVersion\": \"15206257\",
    \"selfLink\": \"/api/v1/namespaces/longhorn-system\",
    \"uid\": \"9673789f-14fb-11e9-ba68-005056b171b1\"
  },
  \"spec\": {
    \"finalizers\": [
      \"kubernetes\"
    ]
  },
  \"status\": {
    \"phase\": \"Terminating\"
  }
}
```

Chapter 16: Setting up a CI/CD pipeline and image registry



```
## Create drone namespace
kubectl create ns drone --dry-run=client | kubectl apply -f

## Setup helm repo
helm repo add bitnami https://charts.bitnami.com/bitnami

## Install the chart
helm install drone-db bitnami/postgresql -n drone \
--set global.storageClass=longhorn \
--set global.postgresql.auth.postgresPassword=drone \
--set global.postgresql.auth.username=drone \
--set global.postgresql.auth.password=drone \
--global.postgresql.auth.database=drone
```

```

## Setup helm repo
helm repo add drone https://charts.drone.io
helm repo update

## Upload ssl cert
kubectl -n drone create secret tls ssl-cert --key="tls.key" --cert="tls.crt"

## Create RPC Secret
openssl rand -hex 16
bea26a2221fd8090ea38720fc445eca6

## Install the chart
helm install drone-server drone/drone -n drone \
--set ingress.enabled=true \
--set ingress.hosts[0].host=drone.example.com \
--set ingress.tls[0].hosts=drone.example.com \
--set ingress.tls[0].secretName=ssl-cert \
--set persistentVolume.enabled=true \
--set persistentVolume.storageClass=longhorn \
--set env.DRONE_SERVER_HOST=drone.example.com \
--set env.DRONE_SERVER_PROTO=https \
--set env.DRONE_DATABASE_DRIVER=postgres \
--set env.DRONE_DATABASE_DATASOURCE="postgres://drone:drone@drone-db:5432/drone?sslmode=disable" \
--set env.DRONE_RPC_SECRET=bea26a2221fd8090ea38720fc445eca6 \
--set env.DRONE_GIT_ALWAYS_AUTH=true \
--set env.DRONE_GITHUB_CLIENT_ID=your-id \
--set env.DRONE_GITHUB_CLIENT_SECRET=github-secret \
--set env.DRONE_USER_CREATE=username:Your-GitHub-Username,admin:true \
--set env.DRONE_USER_FILTER=Your-GitHub-ORG

```

```

## Create drone namespace
kubectl create ns drone-runner --dry-run=client | kubectl apply -f

## Setup helm repo
helm repo add drone https://charts.drone.io
helm repo update

## Install the chart
helm install kube-runner drone/drone -n drone-runner \
--set env.DRONE_SERVER_HOST=drone.example.com \
--set env.DRONE_SERVER_PROTO=https \
--set env.DRONE_RPC_SECRET=bea26a2221fd8090ea38720fc445eca6

```

```

## Create a ServiceAccount and assign it ClusterAdmin permissions
kubectl -n kube-system create serviceaccount drone
kubectl create clusterrolebinding --clusterrole=cluster-admin --serviceaccount=kube-system:drone
TOKENNAME=`kubectl -n kube-system get serviceaccount/drone -o jsonpath='{.secrets[0].name}'`
TOKEN=`kubectl -n kube-system get secret $TOKENNAME -o jsonpath='{.data.token}' | base64 --decode`
echo $TOKEN

```

```
apiVersion: v1
kind: Config
clusters:
- name: "example-cluster"
  cluster:
    server: "https://rancher.example.com/k8s/clusters/c-m-abcdefgj"

users:
- name: "example-cluster"
  user:
    token: "kubeconfig-user-abcdefgj:1234678912346789123467891234678912346789"

contexts:
- name: "example-cluster"
  context:
    user: "example-cluster"
    cluster: "example-cluster"

current-context: "example-cluster"
```

Repositories >

hello-world

Builds Branches Deployments Settings

General

Secrets

Cron Jobs

Badges

ORGANIZATION

Secrets

Templates

+ NEW SECRET

Secrets

parameters, such as
and ssh keys.

Create a New Secret

Name

Value

☐ Allow Pull Requests

Create Cancel

```
docker-username: DockerHub-Username
docker-password: DockerHub-PAT
kubernetes_server: https://rancher.example.com/k8s/clusters/c-m-abcdefgj
kubernetes_token: abcdefghiklmnopqrstuvwxyz1234567890.....
```

```
kind: pipeline
type: kubernetes
name: Example-App

steps:
- name: Docker-Build
  image: plugins/docker
  settings:
    tags:
      - ${DRONE_BUILD_NUMBER}
      - ${DRONE_COMMIT_BRANCH}
      - latest
    username:
      from_secret: docker-username
    password:
      from_secret: docker-password

- name: Deploy-to-k8s
  image: supporttools/kube-builder:latest
  settings:
    kubernetes_server:
      from_secret: kubernetes_server
    kubernetes_token:
      from_secret: kubernetes_token
  commands:
    - bash /usr/local/bin/init-kubectl
    - kubectl create ns hello-world --dry-run=client | kubectl apply -f
    - kubectl apply -n hello-world -f deploy.yaml
```

```
## Setup helm repo
helm repo add harbor https://helm.goharbor.io
helm repo update

## Create a namespace
kubectl create ns harbor --dry-run=client | kubectl apply -f

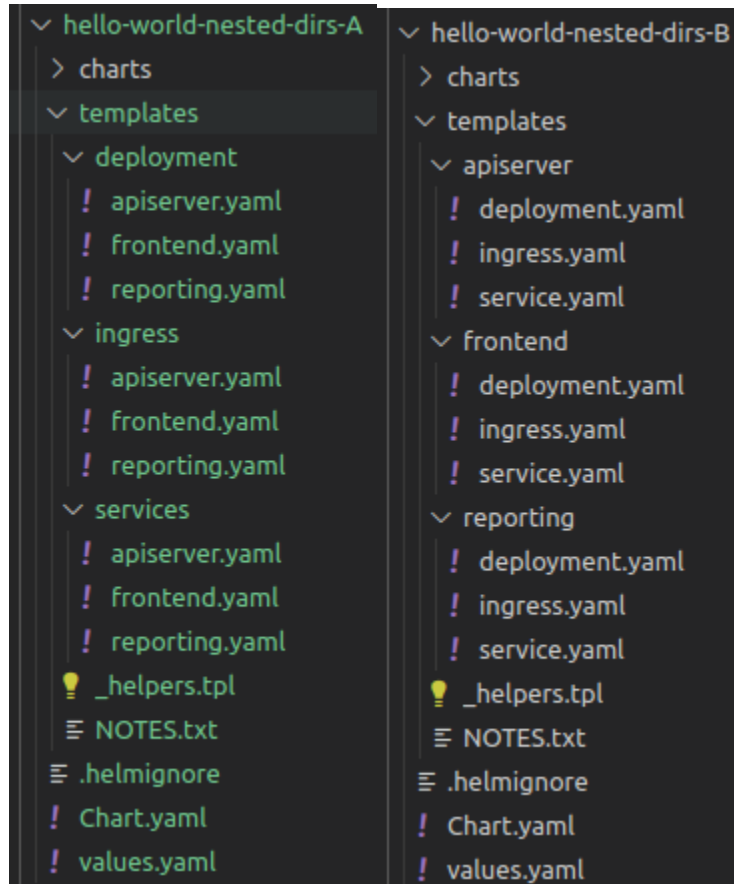
## Upload ssl cert
kubectl -n harbor create secret tls ssl-cert --key="tls.key" --cert="tls.crt"

## Install the chart
helm upgrade --install harbor harbor/harbor -n harbor \
--set expose.type=ingress \
--set expose.tls.secret.secretName=ssl-cert \
--set expose.tls.secret.notarySecretName=ssl-cert \
--set expose.ingress.hosts.core=harbor.example.com \
--set expose.ingress.hosts.notary=harbor-notary.example.com \
--set persistence.enabled=true \
--set persistence.persistentVolumeClaim.registry.storageClass=longhorn \
--set persistence.persistentVolumeClaim.chartmuseum.storageClass=longhorn \
--set persistence.persistentVolumeClaim.jobservice.storageClass=longhorn \
--set persistence.persistentVolumeClaim.database.storageClass=longhorn \
--set persistence.persistentVolumeClaim.redis.storageClass=longhorn \
--set persistence.persistentVolumeClaim.trivy.storageClass=longhorn
```


Chapter 17: Creating and using Helm charts

```
hello-world/  
  charts/          # A directory containing any charts upon which this chart depends.  
  Chart.yaml       # A YAML file containing information about the chart  
  values.yaml      # The default configuration values for this chart  
  templates/       # A directory of templates that, when combined with values,  
                  # will generate valid Kubernetes manifest files.
```

```
apiVersion: v2  
name: hello-world  
description: hello-world  
type: application  
version: 0.1.0  
appVersion: 1.0.0  
keywords:  
  - hello-world  
sources:  
  - https://github.com/SupportTools/hello-world  
maintainers:  
  - name: Matthew Mattox  
    email: mmatttox@support.tools  
    url: https://github.com/mattmatttox/  
icon: https://raw.githubusercontent.com/SupportTools/hello-world/main/img/logo.svg
```



```

frontend:
  image: supporttools/hello-world
  tag: v0.1.2

  ingress:
    enabled: true
    hosts:
      - host: hello-world.example.local
        paths:
          - path: /
            pathType: ImplementationSpecific

apiserver:
  image: supporttools/hello-world-api
  tag: v0.1.1-rc1

  ingress:
    enabled: false
    hosts:
      - host: chart-example.local
        paths:
          - path: /
            pathType: ImplementationSpecific

```

```

{{- if .Values.ingress.enabled }}
{{- if or (.Capabilities.APIVersions.Has "networking.k8s.io/v1/Ingress") (not (.Capabilities.APIVersions.Has "networking.k8s.io/v1beta1/Ingress")) }}
apiVersion: networking.k8s.io/v1
{{- else }}
apiVersion: networking.k8s.io/v1beta1
{{- end }}
kind: Ingress
metadata:

```

```
"Error: UPGRADE FAILED: rendered manifests contain a resource that already exists.  
Unable to continue with update: Deployment "sentry-relay" in namespace "infra-sentry" exists and cannot  
be imported into the current release: invalid ownership metadata; label validation error: missing key  
"app.kubernetes.io/managed-by": must be set to "Helm"; annotation validation error: missing key  
"meta.helm.sh/release-name": must be set to "sentry"; annotation validation error: missing key  
"meta.helm.sh/release-namespace": must be set to "infra-sentry"
```

```
▼ mychart  
  ▼ charts  
  ▼ templates  
    > tests  
    💡 _helpers.tpl  
    ! deployment.yaml  
    ! hpa.yaml  
    ! ingress.yaml  
    ≡ NOTES.txt  
    ! service.yaml  
    ! serviceaccount.yaml  
    ≡ .helmignore  
    ! Chart.yaml  
    ! values.yaml  
    ⓘ README.md
```

```
echo "Deploying Portal"  
helm upgrade --install portal ./chart \  
--namespace ${namespace} \  
-f ./chart/values.yaml \  
--set image.tag=${DRONE_BUILD_NUMBER} \  
--set ingress.host=${ingress}
```

```
- name: Deploy-to-Dev
  image: supporttools/kube-builder:latest
  environment:
    DOCKER_USERNAME:
      from_secret: harbor-username
    DOCKER_PASSWORD:
      from_secret: harbor-password
  settings:
    kubernetes_server:
      from_secret: k8s_dev_server
    kubernetes_token:
      from_secret: k8s_dev_token
  commands:
    - bash deploy.sh dev ${TAG}
```

Chapter 18: Resource management

```
apiVersion: v1
kind: Pod
metadata:
  name: hello-world
spec:
  containers:
  - name: webserver
    image: supporttools/hello-world
    resources:
      requests:
        memory: "64Mi"
        cpu: "250m"
      limits:
        memory: "128Mi"
        cpu: "500m"
```

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: compute-resources
spec:
  hard:
    requests.cpu: "1"
    requests.memory: 1Gi
    limits.cpu: "2"
    limits.memory: 2Gi
    requests.nvidia.com/gpu: 4
    configmaps: "10"
    persistentvolumeclaims: "4"
    pods: "4"
    replicationcontrollers: "20"
    secrets: "10"
    services: "10"
    services.loadbalancers: "2"
```

```
kubectl create namespace kubecost
helm repo add kubecost https://kubecost.github.io/cost-analyzer/
helm install kubecost kubecost/cost-analyzer --namespace kubecost --set kubecostToken="abc123....."
```